

From DEPARTMENT OF CLINICAL NEUROSCIENCE

Karolinska Institutet, Stockholm, Sweden

**COGNITIVE CONTROL IN
REINFORCEMENT LEARNING - BRAIN
STRUCTURE AND FUNCTION**

Björn C. Schiffler



Stockholm 2017

© 2017 Björn C. Schiffler

All previously published papers were reproduced with permission from the publisher

Published by Karolinska Institutet

Printed by US-AB

ISBN 978-91-7676-883-9

Cognitive Control in Reinforcement Learning - Brain Structure and Function

THESIS FOR DOCTORAL DEGREE (Ph.D.)

By

Björn C. Schiffler

Principal Supervisor:

Dr. Sara L. Bengtsson
Karolinska Institutet
Department of Clinical Neuroscience

Opponent:

Dr. Claudia Danielmeier
University of Nottingham
School of Psychology

Co-supervisors:

Prof. Martin Ingvar
Karolinska Institutet
Department of Clinical Neuroscience

Prof. Peter Fransson
Karolinska Institutet
Department of Clinical Neuroscience

Examination Board:

Prof. Martin Lövdén
Karolinska Institutet
Department of Neurobiology, Care Sciences and Society

Prof. Peter Juslin
Uppsala Universitet
Department of Psychology

Dr. Marc Guitart-Masip

Karolinska Institutet
Department of Neurobiology, Care Sciences and Society

Abstract

Adapting behaviour according to internal or external feedback is a fundamental property of cognitive control. For example, humans tend to slow down when they make mistakes, a process called post-error slowing (PES), which has previously received extensive attention in research on response inhibition. However, whether PES is actually an adaptive process which helps avoid future mistakes or a maladaptive one which siphons cognitive resources is still not clear.

The overall aim of the work in this thesis was to investigate how post-error slowing contributes to the stabilization of performance after errors and which brain areas are involved in this response inhibition process. We did this by combining behavioural experiments, computational modelling and neuroimaging techniques to provide a comprehensive analysis of latent decision processes and their neural correlates.

Specifically, in **Study I**, we analyzed data from a probabilistic reinforcement learning task in combination with functional Magnetic Resonance Imaging to explore which brain regions signalled enhanced future post-error slowing when receiving negative feedback. On a behavioural level, we studied whether PES was associated with how well participants learned, as assessed in a later test phase. We showed that post-error slowing was associated with brain activity in a central cognitive control region, the right inferior frontal gyrus (rIFG) as well as brain regions in occipital cortex which overlapped with the representation of absolute prediction errors, a measure reflecting deviance from expectations, i.e., surprise at feedback. In **Study II**, we found that cortical thickness in rIFG as a measure of grey matter integrity was related to inter-individual differences in post-error slowing, both for direct next trials and trials further apart in time. This analysis was supported by a drift diffusion model of the underlying decision processes, which demonstrated that an increased decision boundary after an error, indicating enhanced response caution, was related to cortical thickness variability,

particularly in anterior parts of the rIFG. Finally, in **Study III** we used drift diffusion modelling on a large-scale behavioural dataset during a visual search task to illuminate decision processes of up to five trials after an error and how post-error adaptation benefits accuracy recovery several trials after the error. Post-error slowing was marked by both adaptive and non-adaptive decision processes which changed dynamically over several trials after an error. While adaptive increases in decision threshold were sustained for several trials after an error, reductions in evidence accumulation only transiently affected performance on the next trial after the error. Further, post-error increases in response caution and evidence accumulation were also associated with better performance on future trials.

These studies illustrate that there is valuable information to be gained about response inhibition processes beyond looking at the simple relation of post-error slowing and accuracy. Computational modelling allowed us to compartmentalize various decision processes and relate adaptations in these processes directly to brain anatomy. We hope the results from the studies presented in this thesis can provide a framework for future work on how the brain learns from mistakes and adapts to a continually changing environment.

List of publications

Schiffler BC, Almeida R, Granqvist M, Bengtsson SL (2016). Memory-reliant post-error slowing is associated with successful learning and fronto-occipital activity. *Journal of Cognitive Neuroscience*, 28, 1539-1552.

Schiffler BC, Abé C, & Bengtsson SL (2017). Cortical thickness of right inferior frontal cortex is associated with memory-based post-error slowing. *Manuscript*.

Schiffler BC, Bengtsson SL[†], & Lundqvist D[†] (2017). The Sustained Influence of an Error on Future Decision-Making. *Frontiers in Psychology*, 8, 1077. [†]Equal contribution.

List of additional publications

Plavén-Sigray P*, Matheson GJ*, **Schiffler BC***, Thompson WH (2017). The readability of scientific texts is decreasing over time. *eLife*, 6, e27725. *Equal contribution.

Contents

Abbreviations	9
1 Introduction	10
1.1 Conceptual overview	10
1.2 Cognitive control	11
1.3 Reinforcement learning	15
1.4 The intersection of cognitive control and reinforcement learning .	21
1.5 Challenges	22
1.6 Computational modelling to tackle these challenges	23
2 Aims	29
3 Methodological Considerations	31
3.1 Magnetic Resonance Imaging	31
3.2 Functional Magnetic Resonance Imaging	33
3.3 Reinforcement Learning Modelling	34
3.4 Hierarchical Bayesian Estimation of the Drift Diffusion Model (HDDM)	35
3.5 Post-error slowing	35
3.6 Participants	37
4 Results	38

4.1	Study I: Post-error slowing is associated with learning performance and functional activity in cognitive control and visual regions	38
4.2	Study II: Adaptive increases in response caution after errors are related to cortical thickness in cognitive control regions	42
4.3	Study III: Response adaptations to errors are multi-component processes and change dynamically over several trials after the error	44
5	Discussion	46
5.1	Behavioural results in Studies I - III	46
5.2	Neuroimaging results in Studies I and II	48
5.3	Limitations and future directions	49
	Acknowledgments	53
	References	55

Abbreviations

Abbreviation	Term
ACC	Anterior Cingulate Cortex
BOLD	Blood Oxygen Level Dependent
DDM	Drift Diffusion Model
fMRI	Functional Magnetic Resonance Imaging
MRI	Magnetic Resonance Imaging
PES	Post-Error Slowing
pMFC	Posterior Medial Frontal Cortex
PFC	Prefrontal Cortex
pre-SMA	Presupplementary Motor Area
RT	Reaction Time
ROI	Region Of Interest
RL	Reinforcement Learning
rIFC	Right Inferior Frontal Cortex
rIFG	Right Inferior Frontal Gyrus
SSRT	Stop Signal Reaction Time
STN	Subthalamic Nucleus
2AFC	Two-Alternative Forced Choice

Chapter 1

Introduction

One of the lines of experimental investigation most diligently followed of late years is that of the ascertainment of the time occupied by nervous events. [...] The question is, What happens inside of us, either in brain or in mind? and to answer that we must analyze just what processes the reaction involves.

– William James, *Principles of Psychology*

1.1 Conceptual overview

Both on the behavioural and neural level tremendous progress has been made over the past decades in the understanding of how stimulus values lead to action and how actions can be constrained by exerting cognitive control over them. However, these two processes of reinforcement learning and cognitive control have predominantly been investigated in isolation, particularly in neuroscience. Therefore, the question of how appropriate behavioural adjustment according to feedback improves value

learning has only recently come into the spotlight. This question will stand in the center of this thesis.

I will start by providing separate overviews of the fields of cognitive control and reinforcement learning and afterwards I am going to highlight interactions between the two fields, particularly in the realm of mental disorders.

1.2 Cognitive control

Cognitive control refers to a heterogeneous concept, subsuming a variety of mental functions which include, but are not limited to, working memory, task-set switching, response selection and response inhibition (Lenartowicz, Kalar, Congdon, & Poldrack, 2010; Ridderinkhof, Widenberg, Segalowitz, & Carter, 2004; Sabb et al., 2008; Ullsperger, Danielmeier, & Jocham, 2014). As a whole it can be defined as “the ability to regulate, coordinate, and sequence thoughts and actions in accordance with internally maintained behavioral goals” (Braver, 2012). In this thesis I will mainly focus on behavioural and neural aspects of response inhibition in order to implement adaptive behaviour.

1.2.1 Post-error slowing (PES)

An example of response inhibition is the phenomenon of post-error slowing (PES). After making a mistake or receiving negative feedback in a task, the next response on similar trials will often require a longer reaction time than pre-error responses (Danielmeier & Ullsperger, 2011; Kerns et al., 2004).

Various accounts have been proposed to explain the psychological correlates to this slowing. On the one hand, non-functional accounts of PES contend that the slowing

happens because of the low frequency of errors in many tasks which evokes a general orienting response (Notebaert et al., 2009) or that post-error monitoring acts as a bottleneck, limiting resources of decision processes after the error (Jentzsch & Dudschig, 2009). A common motive among the functional accounts of PES on the other hand is to argue for an increase in response caution (Botvinick, Braver, Barch, Carter, & Cohen, 2001; Dutilh et al., 2012b) on the trial following the error which leads to a speed-accuracy trade-off (Bogacz, Wagenmakers, Forstmann, & Nieuwenhuis, 2010; Fitts, 1966; Heitz, 2014; Laming, 1979; Steinhauser, Ernst, & Ibald, 2017).

PES is often quantified in response time tasks such as Go/No-Go tasks, where the response time after an error is subtracted from the response time on the trial before the error (Dutilh et al., 2012a). Similarly, it can be calculated in relation to a comparable previous trial, even if that trial is set apart in time by several seconds (Cavanagh, Frank, Klein, & Allen, 2010; Frank, Moustafa, Haughey, Curran, & Hutchison, 2007). While the former might reflect a motoric or more general attentional adjustment, the latter also takes into account the adaptation of response speed in accordance to conflict between two stimuli with similar values.

For example, in a probabilistic reinforcement learning task such as we have employed in **Study I** and **II**, negative feedback should trigger response speed decreases the next time a particular symbol pair is seen, which would show an adaptive regulation of response time in accordance with feedback (Figure 1.1A). This response time slowing could be useful to gather more evidence when stimulus values are more similar to each other and therefore harder to distinguish, i.e., during conflict (Cavanagh et al., 2011; Frank et al., 2015; Zaghloul et al., 2012). On the other hand, errors on a visual search task such as the one we used in **Study III** might be expected to have the largest effect on trials directly after the error since the task is structured deterministically and errors provide information about necessary adaptation on the immediate next trial (e.g., “I missed this angry

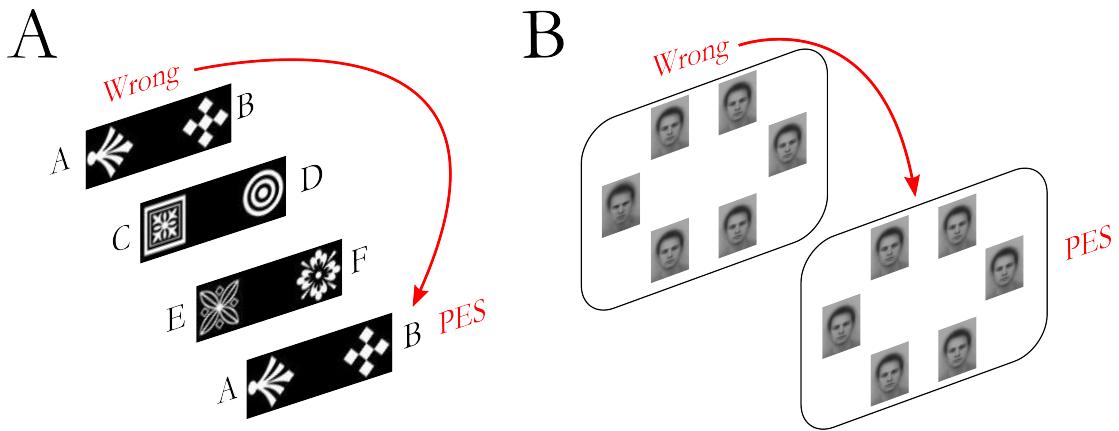


Figure 1.1: The two different kinds of post-error slowing investigated in this thesis. (A) PES during a probabilistic reinforcement learning task as employed in our **Studies I** and **II**. Here, PES (ΔRT) is calculated by subtracting the RT on the error trial from the RT when the same pair trial next appears again. (B) PES during a visual search task as used in our **Study III**. In this study, ΔRT is calculated by subtracting the RT on the trial after the error from the RT on the trial before the error if both trial conditions are the same (in this case all neutral faces).

face, therefore I need to check for this particular face shape more carefully”, see Figure 1.1B).

1.2.2 The cognitive control network in the brain

Cognitive control processes are implicated to rely on lateral prefrontal brain areas (see e.g., the review by Miller, 2000) and anterior cingulate cortex as indicated in a lesion mapping review with a variety of neuropsychological tasks related to response inhibition, conflict monitoring and switching (Gläscher et al., 2012). A particular network that implements response inhibition or stopping depending on the task at hand has been proposed by Aron and colleagues (Figure 1.2A, Aron, Behrens, Smith, Frank, & Poldrack, 2007; Aron, Robbins, & Poldrack, 2014). Within this network, right inferior frontal gyrus (rIFG) plays a key role in response inhibition via its connections to basal ganglia (subthalamic nucleus, STN) and

presupplementary motor area (pre-SMA). Anatomical connections between the areas in this cognitive control network show good test-retest reliability as assessed in a recent diffusion weighted imaging study (Boekel, Forstmann, & Keuken, 2017).

PES has been linked to posterior medial frontal cortex (pMFC) activity, increased activity in task-relevant areas and decreased activity in task-irrelevant cortical areas (Danielmeier, Eichele, Forstmann, Tittgemeyer, & Ullsperger, 2011; J. A. King, Korb, Cramon, & Ullsperger, 2010). The interesting dynamic here is that higher activity in pMFC after feedback was correlated with an increase in PES on the post-error trial (Danielmeier et al., 2011) while decreases in motor network activity on the post-error trial were related to a decrease in PES (J. A. King et al., 2010).

A recent study combining effective and anatomical connectivity methods (Dynamic Causal Modelling and diffusion based probabilistic tractography) showed that the rIFG positively modulates the excitatory influence of the pre-SMA on the STN, leading to stronger motor inhibition in motor cortex (Figure 1.2B, Rae, Hughes, Anderson, & Rowe, 2015). Importantly, mean diffusivity in white matter tracts connecting rIFG and preSMA to STN correlated with better performance in inhibiting responses in a stop-signal task.

Extent of lesion damage to rIFG has previously been found to correlate with a classical measure of response inhibition, the stop signal reaction time (SSRT) in a stop signal task (Aron, Fletcher, Bullmore, Sahakian, & Robbins, 2003) and temporary deactivation of the rIFG using transcranial magnetic stimulation lead to the inability of stopping an action (Chambers et al., 2006).

Further, studies have shown that white matter integrity connecting the rIFG and other areas relevant for cognitive control showed an association to response inhibition performance (Fjell, Westlye, Amlie, & Walhovd, 2012; Madsen et al., 2010).

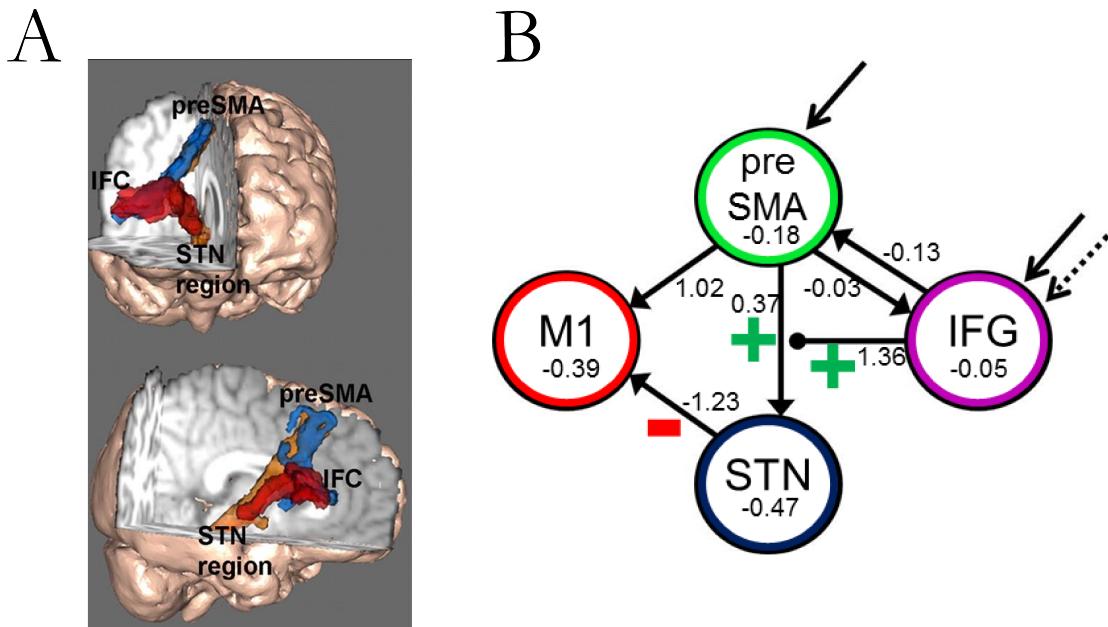


Figure 1.2: Major nodes of a prefrontal cognitive control network. (A) White matter tracts connecting the right IFC, preSMA and STN are visualized using Diffusion Weighted Imaging. Figure from Aron et al. (2007), reprinted with permission from Society for Neuroscience. (B) Dynamic Causal Modelling of the cognitive control network. The winning model configuration suggests a modulating influence of IFG on the activity between preSMA and STN. Figure from Rae et al. (2015), reprinted under CC BY 4.0 license.

Based on these results of a crucial role of rIFG in response inhibition, in **Study II** we therefore evaluated the hypothesis whether cortical thickness in rIFG could be related to decision components involved in post-error slowing.

1.3 Reinforcement learning

One of the largest influences on modern approaches to reinforcement learning are behavioural studies in animals, carried out at the beginning of the last century. In the framework of classical conditioning, Pavlov used the term “reinforcement” as the repeated pairing between an unconditioned (e.g., food) and a neutral stimulus (e.g., a bell) until the neutral stimulus by itself led to the response (e.g., dogs’

salivation) previously only associated with the unconditioned stimulus (Pavlov, 1927). Thus, reinforcement here refers to an association between two stimuli.

Thorndike on the other hand specified an action-stimulus association in his Law of Effect, derived from seminal experiments with cats (Thorndike, 1911). This law pertains to the observation that behaviour that is followed by positive consequences is likely going to be repeated in the future, e.g., a cat pressing a lever to escape a box and obtain the fish outside the box, a process later on referred to as instrumental or operant conditioning (Skinner, 1935). Modern computational approaches to reinforcement learning implement both Pavlovian and instrumental conditioning to model behaviour by utilizing algorithms from the field of machine learning (Sutton & Barto, 1998).

A highly influential model depicting Pavlovian conditioning in animals was developed by Rescorla & Wagner (1972). This model was particularly popular as it could explain key behavioural findings in reinforcement learning such as blocking (Niv, 2009). Blocking refers to the finding that a second conditioned stimulus does not evoke a conditioned response if it does not provide additional information beyond that of the first conditioned stimulus (Kamin, 1969).

Sutton and Barto (Sutton, 1988; Sutton & Barto, 1990) extended these ideas with the concept of temporal difference learning, incorporating future rewards, discounted by how far they were set apart in time. This can be a useful property when trying to explain processes like second-order conditioning and conditioned reinforcement (Niv & Schoenbaum, 2008). In second-order conditioning, a stimulus that had previously been conditioned can then be associated with another stimulus to construct a chain of associations. For example, an animal which has learned to associate the ringing of a bell with food can be taught to associate a light with the bell alone and through those contingencies learn a higher-order association between light and food. Temporal difference learning has been successfully applied

to describe this type of second-order conditioning in humans (Seymour et al., 2004).

Here, the notion of a prediction error was also introduced, indicating the calculation of a predicted reward minus the actual obtained reward. I will refer back to the prediction error, particularly in terms of reward learning, in the section on neural correlates of reinforcement learning.

Actions that lead to rewards are subsequently repeated. This is taken into account in Q-learning models (Watkins, 1989; Watkins & Dayan, 1992). Here, state-action pairs, instead of state-value pairs, are modelled as Q-values with the “best” actions referring to the ones with the highest Q-values (Niv, 2009). The estimated values can then also be used to calculate a prediction error. The SARSA (state-action-reward-state-action) algorithm on the other hand is considered an on-policy algorithm, meaning that the calculation of the prediction error only involves the next chosen action, i.e., following the agent’s policy (Niv, 2009). In recent years, both using Q-learning (e.g., Fitzgerald, Friston, & Dolan, 2012; Frank et al., 2007) and SARSA (e.g., Daw, 2011; Gläscher, Daw, Dayan, & O’Doherty, 2010) algorithms to model human behaviour have proved to be useful and reinforcement learning approaches have also improved performance of artificial intelligence algorithms such as deep learning to match or improve upon human performance (Mnih et al., 2015; Silver et al., 2017).

Furthermore, a division between model-free and model-based reinforcement learning has also been proposed. While the former refers to the simple storing of action-value correspondences without assuming a causal structure of the environment, the latter emphasizes that decisions which lead to rewards are made through planning and simulating the potential actions in a goal-directed manner (Botvinick & Weinstein, 2014). Whether these two different systems are also dissociable in terms of their neural correlates is still being discussed, see e.g., Daw (2011) for evidence in

favour of an integrated system, Gläscher et al. (2010) and Smittenaar, FitzGerald, Romei, Wright, & Dolan (2013) for accounts of separable neural regions, S. W. Lee, Shimojo, & O'Doherty (2014) for an arbitration mechanism between the two systems, as well as Wunderlich, Smittenaar, & Dolan (2012) for the influence of dopamine in that context.

1.3.1 Neural correlates of reinforcement learning

In their application to neuroscience, a crucial assumption underlies many of these reinforcement learning models, namely the idea that the brain continually makes predictions about its environment in order to maximize reward (Cohen, McClure, & Yu, 2007; see also Friston & Kiebel, 2009; Niv, 2009).

One of the most intriguing findings bringing together computational approaches and neural processes was the discovery that neurons expressing the neurotransmitter dopamine in nonhuman primates not only reacted to rewarding stimuli but also shifted their response towards a cue that could reliably predict an upcoming reward while not firing when the actual reward was presented. Furthermore, a decrease in firing rate was observed when the original reward was omitted at the expected point in time. In other words, the phasic firing of these neurons seemed to achieve a reward prediction error (Montague, Dayan, & Sejnowski, 1996; Schultz, Dayan, & Montague, 1997, see Figure 1.3).

While not uncontroversial (Redgrave, Gurney, & Reynolds, 2008; Redgrave, Prescott, & Gurney, 1999), this finding of a direct neural implementation of a prediction error and the close association between dopamine and reinforcement learning has inspired a lot of research in human neuroimaging, commonly attributing neural correlates of reward prediction errors to the striatum (see e.g., Chase, Kumar, Eickhoff, & Dombrovski, 2015; Garrison, Erdeniz, & Done, 2013

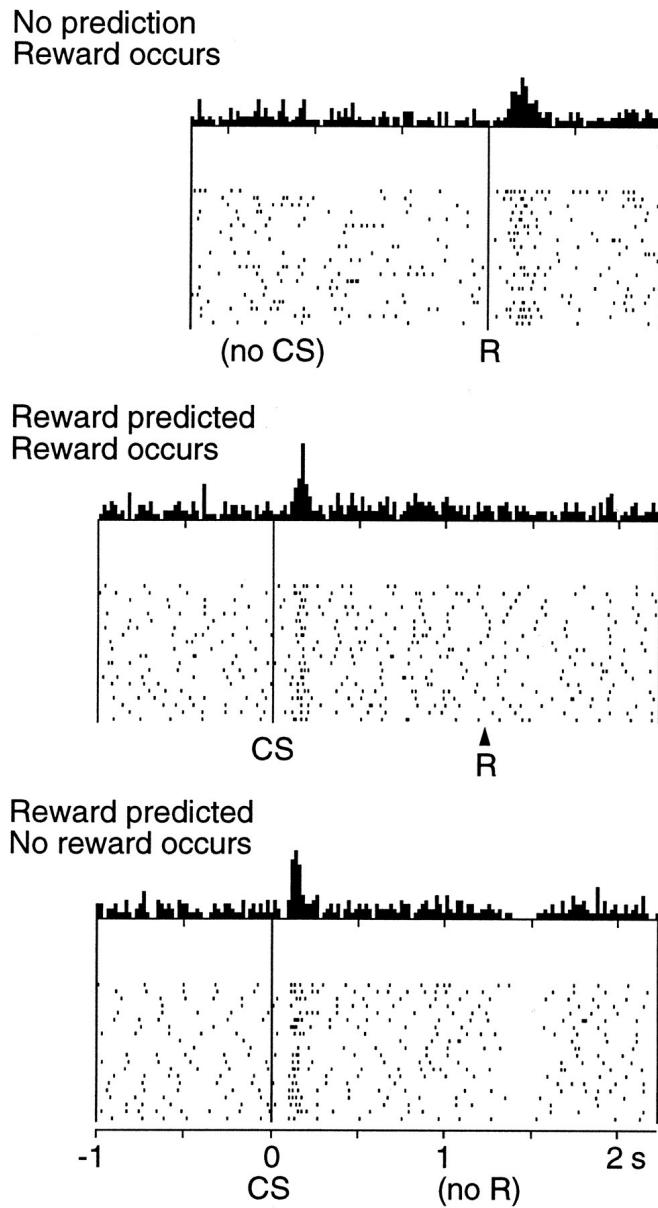


Figure 1.3: Dopamine neurons implement a reward prediction error. When no conditioned stimulus (CS) is presented, dopamine neurons fire in response to a reward R (top panel). When a CS such as a light is predictive of the reward, the dopamine response shifts to firing after the CS instead (middle panel). If no reward is presented even though it was predicted because of the light, there is a dip in the dopaminergic neuron response at the time when the reward should have appeared. Figure from Schultz et al. (1997), reprinted with permission from AAAS.

for recent reviews).

Several studies have shown that dopamine levels in the brain promote associative learning in specific ways. Frank and colleagues, using a probabilistic learning task, discovered that dopamine medication in Parkinson's patients would make them better at learning from positive than negative outcomes, while patients off medication showed the opposite contrast (Frank, Seeberger, & O'Reilly, 2004). Further, polymorphisms in genes related to dopamine function also showed a direct relation to positive and negative learning styles. For example, participants with variations in dopamine D2 receptor densities displayed differences in their proneness to learn from or generalize to negative outcomes (Frank et al., 2007; T. A. Klein et al., 2007).

Computational modelling has been effectively used in combination with neuroimaging methods like fMRI to elucidate e.g., developmental changes in reinforcement learning (Van Den Bos, Cohen, Kahnt, & Crone, 2012), risk sensitivity (Niv, Edlund, Dayan, & O'Doherty, 2012) and deviations in reward learning in mental disorders like depression (Kumar et al., 2008) and bulimia nervosa (G. K. W. Frank, Reynolds, Shott, & O'Reilly, 2011). Neuroimaging studies provided additional evidence that striatum and midbrain are involved in the computation of reward prediction errors and implicate the ventromedial prefrontal cortex in the estimation of values (Chase et al., 2015; Jocham, Klein, & Ullsperger, 2011).

1.4 The intersection of cognitive control and reinforcement learning

How cognitive control processes can benefit value learning, override action tendencies or sway neural processes to promote a more model-based or model-free computation has only recently sparked investigations. For example, in a recent study, cognitive control related measures in two tasks were associated with a higher amount of model-based reward learning in a two-stage reinforcement learning task (Otto, Skatova, Madlon-Kay, & Daw, 2015).

Yet, whether cognitive control can also have an impact on model-free reinforcement learning is less well known. We have addressed this issue in **Study I**, demonstrating that post-error slowing can have an impact on the generalization of learning in a model-free reinforcement learning paradigm (Schiffner, Almeida, Granqvist, & Bengtsson, 2016). Another study provided support that connectivity between anterior cingulate cortex and ventromedial prefrontal cortex was associated with adaptive switches in choice behaviour, integrating both immediate and delayed consequences (Economides, Guitart-Masip, Kurth-Nelson, & Dolan, 2014). This finding suggests a possible cognitive control process, foregoing the immediate reward for a potentially larger delayed outcome.

Furthermore, the interaction between cognitive control and reward learning plays an important role in mental disorders and can thus be a potential future target for specific intervention. Particularly in addiction, self-control over prepotent response habits is required to sustain abstinence. Volkow and colleagues showed that instructed cognitive control of drug craving activated right inferior frontal gyrus as a node in a cognitive control network while inhibiting reward related regions, including nucleus accumbens and orbitofrontal cortex (Volkow et al., 2010). As compulsive disorders show a habit towards model-free at the expense of

model-based learning (Voon et al., 2015), further work on the relation of cognitive control and reward learning styles could be beneficial to illuminate the ways in which cognitive control can benefit abstinence and ultimately recovery from these disorders.

1.5 Challenges

1.5.1 Is response adaptation like PES beneficial to learning?

Whether PES provides specific benefits in acquisition or transfer in learning paradigms is of yet unclear (Ullsperger & Danielmeier, 2016; Ullsperger et al., 2014). For example, it is still to be determined under which conditions PES is associated with post-error accuracy (Danielmeier & Ullsperger, 2011; Hajcak, McDonald, & Simons, 2003; Hester, Barre, Mattingley, Foxe, & Garavan, 2007).

We address these questions in our **Studies I** and **III** (Schiffner et al., 2016; Schiffner, Bengtsson, & Lundqvist, 2017), investigating learning benefits of PES and how post-error decision components relate to stabilisation of accuracy.

1.5.2 How long do PES effects persist after an error?

A related question concerns for how long an error affects decision making. In **Study I**, we investigated whether there was a memory component to PES in a reinforcement learning task, i.e., whether participants would adapt their response speed in relation to the negative feedback on certain symbol pairs. In **Study II**, we teased apart the effects of negative feedback on immediate next trials and delayed adaptation on later trials with regards to anatomical correlates. Finally, in

Study III we investigated the persistence of post-error adaptations several trials after the error in a visual search task.

1.5.3 What are the time courses and contributions to accuracy of decision processes involved in PES?

Previous studies have found to varying degree that both increases in decision threshold and reduced sensitivity to sensory information underlie PES. In **Study III**, we explored how these altered decision components change over time and how they contribute to the stabilisation of accuracy.

1.6 Computational modelling to tackle these challenges

Advances in computer science, neuroscience and artificial intelligence prompted a cognitive revolution in the 1950s which enabled psychology to look beyond the black box of mental phenomena previously favoured by the predominant field of behaviourism. One way to validate theories about mental processes is to utilize computational models which can be fit to behavioural and/or neural data acquired in experiments.

In this section I will present a general overview over the two main models used in our studies. These are reinforcement learning models, which try to explain how participants learn by estimating the stored value of task items and predicting decisions based on these values and drift diffusion models, which are fit to reaction times and accuracy in a given task to illuminate the underlying decision process.

1.6.1 Reinforcement learning modelling

The central theme in reinforcement learning modelling questions how the value of rewards in the environment impacts decision making. The Rescorla-Wagner model of animal learning formalized how the value of a conditioned stimulus (CS) changes in a classical conditioning paradigm dependent on the worth of the actual unconditioned stimulus (US) (formula modified from Niv, 2009; Rescorla & Wagner, 1972):

$$V_{new}(CS_i) = V_{old}(CS_i) + \alpha \left[\lambda_{US} - \sum_i V_{old}(CS_i) \right]$$

Here, learning of the association between US and CS happens because the new value $V_{new}(CS_1)$ of any of the conditioned stimuli (e.g., a light as CS_1) gets updated by adding the difference between what actually happened (e.g., a food pellet as λ_{US}) and what was predicted - $\sum_i V_{old}(CS_i)$ - to its old value $V_{old}(CS_1)$.

The importance of the difference between expected and predicted value was a remarkable insight at the time and is one of the key components of reinforcement learning models and other computational models of brain function even today. It indicates the surprise of a particular outcome and is additionally modulated by a learning rate α which reflects the importance of recency of rewards and can for example vary depending on the salience of the stimuli involved (Niv, 2009).

This measure of surprise also found its way into reinforcement learning models which are in active use in research today. In cognitive neuroscience, the measure of the prediction error has received particular attention because of its close relation with dopamine neuron firing (Schultz et al., 1997) as well as its correspondence to dissociable model-free (i.e., reward-related) and model-based (i.e., state-related) correlates in the brain (Gläscher et al., 2010). Analogous to the final term in

the Rescorla-Wagner model, the prediction error can be calculated as follows by subtracting the expected value of a stimulus V at timestep t from the current value r at timestep t :

$$\delta_t = r_t - V_t$$

The full equation with which value updates can be estimated then looks like this:

$$V_{t+1} = V_t + \alpha * \delta_t$$

How do these stimulus values lead to a decision towards one or another option? The computed values can be converted to action probabilities by the softmax equation, in which for example the probability of choosing between stimulus A and B at time t can be described as follows:

$$P(A)_t = \frac{1}{1 + e^{-\beta(V(A)_t - V(B)_t)}}$$

In this formula, the difference between the values of stimuli A ($V(A)_t$) and B ($V(B)_t$) is computed and adjusted by the inverse temperature parameter β , which controls individuals' propensity to explore new options versus exploit the known value differences (albeit small as they might be). In other words, an increase in the value for β suggests that a participant follows the computed value difference more deterministically.

RL models have already been effectively applied to highlight previously unexplained aspects of disorders (for a review see Maia & Frank, 2011) such as Parkinson's disease (Frank et al., 2004), Schizophrenia (Roiser et al., 2009) and in explaining model-based versus model-free behaviour (Doll, Simon, & Daw, 2012) which might

be of future use to explain pathologies.

1.6.2 Drift diffusion modelling

The fundamental question how mental computations can be inferred from measured reaction times dates back almost 150 years. Already in 1870 and based on earlier computations on nerve conduction velocity by von Helmholtz (1850), Foster described the principle on which today's drift diffusion models (DDM) base their computations, namely the division of overt reaction time into the unobservable components of sensory acquisition, mental computation, and motor output: "A typical bodily action, involving mental effort, may be regarded as made up of three terms ; of sensations travelling towards the brain, of processes thereby set up within the brain, and of resultant motor impulses travelling from the brain towards the muscles which are about to be used" (Foster, 1870).

Modern sequential sampling models, of which DDMs are one subclass, rely on the idea that during speeded decision-making tasks, evidence towards one of the options is acquired in a noisy manner until it reaches a decision boundary, after which a motor action is being executed (Forstmann, Ratcliff, & Wagenmakers, 2016). DDMs have usually been applied to two-alternative forced choice tasks (2AFC), although extensions to include tasks with multiple alternatives have also been discussed (Forstmann et al., 2016; Ratcliff, Smith, Brown, & McKoon, 2016).

Prominent parameters in these models include the decision threshold a , the rate of evidence accumulation v , and the non-decision time T_{er} , which consists of time needed for both sensory acquisition as well as motor execution (see Figure 1.4 for a graphical overview of the DDM and its core parameters). While a reflects how much evidence is necessary to make a decision towards one of the options presented, v indicates the slope at which this evidence is acquired.

These three parameters relate to RT and accuracy in a particular fashion (Table 1.1). The decision threshold a corresponds to increases in response caution with an increase in both RT and accuracy. The evidence accumulation v is related to decreases in RT and simultaneous increases in accuracy and is sensitive to for example task difficulty (higher difficulty leads to decreased v). Finally, the non-decision time T_{er} corresponds to an increase in RT with no concurrent change in accuracy and is hypothesized to be for example affected by aging (Ratcliff, Thapar, & McKoon, 2010), although Ratcliff et al. (2010) also find increasing decision thresholds with advancing age.

Table 1.1: Core parameters of drift diffusion models and their relation to RT, accuracy, and proposed correlates.

DDM Parameter	RT	Accuracy	Proposed correlates
Decision threshold a	↑	↑	Response caution
Drift rate v	↓	↑	Task difficulty
Non-decision time T_{er}	↑	=	Aging

Using reaction time distributions and response accuracy in combination allows for example to explain differences in experimental conditions (Forstmann et al., 2016 for reviews; see Ratcliff et al., 2016) or between patient groups, such as in schizophrenia (A. A. Moustafa et al., 2015), Parkinson’s disease (Zhang et al., 2016) or in patients with psychosis (Mathias et al., 2017) to particular mental processes within the total reaction time.

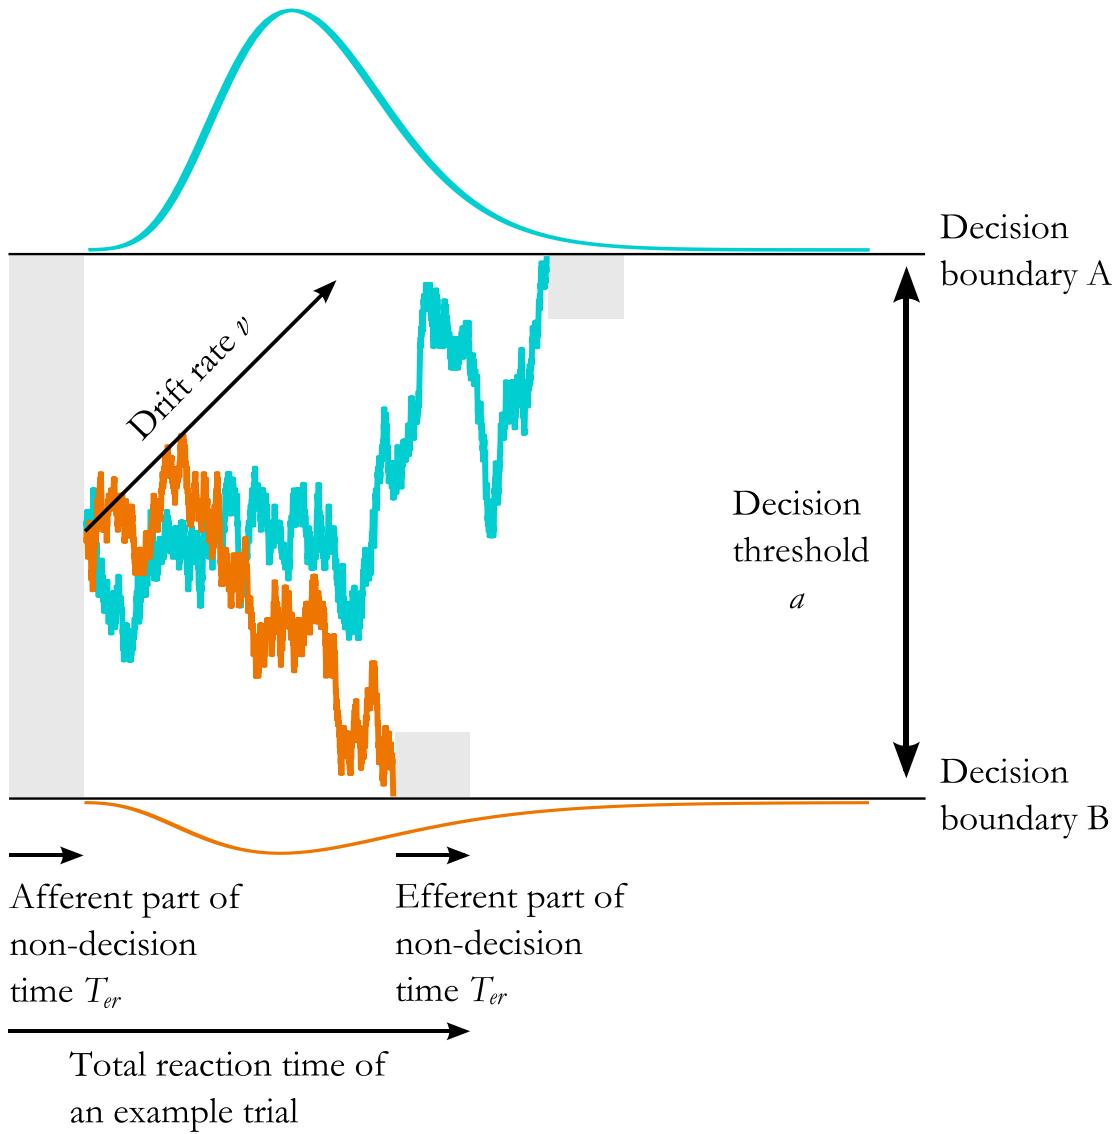


Figure 1.4: Main parameters of the canonical drift diffusion model for decisions in two-alternative forced-choice tasks. Two example diffusion processes as random walks are depicted alongside the three main parameters of the drift diffusion model. Evidence for one or the other option is accumulated with a drift-rate v until one of the two decision boundaries which are separated by the decision threshold a is crossed. The total non-decision time T_{er} can be divided in an afferent part (sensory perception) and an efferent part (movement initiation and execution).

Chapter 2

Aims

The central aim of the studies presented in this thesis was to investigate which aspects of cognitive control processes in reaction to negative feedback benefit learning and how this is reflected in both brain anatomy and function.

In **Study I** we used a probabilistic reinforcement learning paradigm to assess the influence of cognitive control adjustments such as post-error slowing and stay/switch-behaviour in addition to learning phase performance as predictors for learning outcome in a later test phase. In addition, we were interested in which neural areas predicted later response time adaptation at the time of receiving negative feedback. Further, we analyzed how trial-by-trial absolute and signed prediction errors obtained from our reinforcement learning models affected behaviour and how they were represented in the brain. Data for this study was acquired in the context of a larger ongoing project which investigates the influence of self-associations on learning using priming techniques. Therefore we controlled for this factor in all analyses of **Studies I** and **II** in which this was possible.

Given converging evidence from previous research (e.g., Aron et al., 2007, 2014; Rae et al., 2015) and the interesting results from **Study I** which implicated the

right inferior frontal cortex as an important region in cognitive control processes, we investigated in **Study II** whether an anatomical property of this area, cortical thickness, was related to the extent and memory aspects of PES. We also used drift diffusion modelling to better understand the dominant cognitive process behind post-error slowing and related obtained parameters to anatomical structure of the rIFC.

Using a visual search task, in **Study III** we explored the effect of an error on not only the first trial after the error, but also on subsequent trials. This was again supported by drift diffusion models to reveal the relevant decision process components behind reaction times and accuracy. Additionally, we examined how trial type properties (emotion and difficulty) influenced post-error adaptations and how later increases in accuracy were associated with decision processes on the trial after an error.

Chapter 3

Methodological Considerations

In this chapter I will provide an outline of the main techniques used in our studies with a focus on neuroimaging techniques and computational modelling.

3.1 Magnetic Resonance Imaging

Magnetic Resonance Imaging (MRI) is a non-invasive imaging technique that utilizes the fact that different tissue types have different magnetic properties to provide high-resolution images of anatomy such as the brain. To make this possible, a strong magnetic field, commonly in the range of 1.5 - 7 Tesla for human brain imaging, is created by an MRI scanner, which aligns the nuclei of (e.g., hydrogen) atoms in a common direction. The transient introduction of a varying electromagnetic field (B_1) by a radiofrequency pulse leads to excitation and a dephasing of the nuclei. Now it becomes for example possible to measure how long it takes the nuclei to return back to their original state (the so-called *spin-lattice relaxation time* or T_1 relaxation). This property is used in neuroimaging to provide the so-called T_1 contrast, which enables a detailed visualization of differences

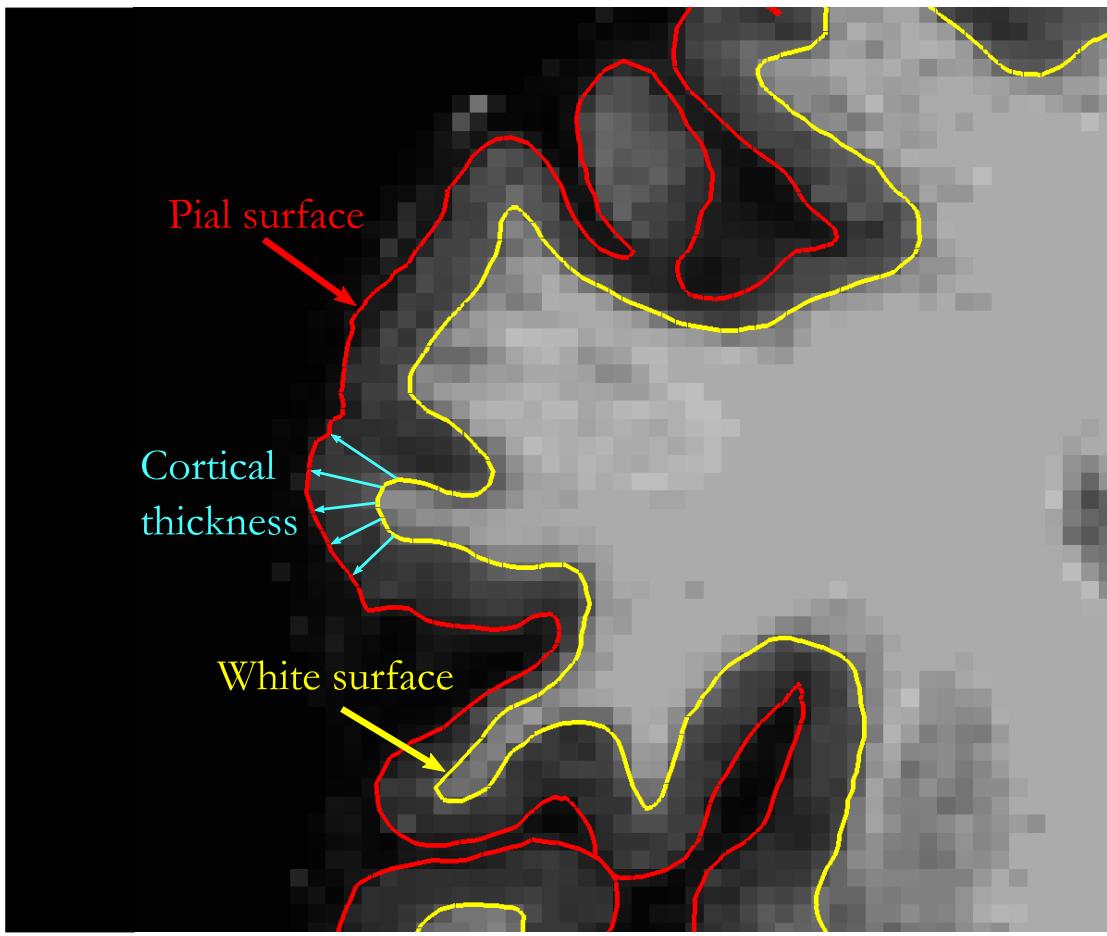


Figure 3.1: Cortical thickness is calculated using the shortest distance between pial surface and white surface. Image was created with the tool TkMedit after running the FreeSurfer standard processing pipeline on one example brain from Study II.

between e.g., grey and white matter in the brain. We make use of the T_1 contrast in **Study II** to assess inter-subject differences in cortical thickness depending on individual propensity to adapt response speed after an error.

Cortical thickness is a MRI-derived metric which has been successfully used as a measure of the integrity of the cerebral cortex (e.g., Dickerson & Wolk, 2012; Makris et al., 2007). It decreases with normal aging (Storsve et al., 2014) and can be related to symptom severity with corresponding regionally specific rates of decrease in Alzheimer's disease (Dickerson et al., 2009). Using freely available

tools such as FreeSurfer (Fischl & Dale, 2000), cortical thickness can be calculated by taking the smallest distance between the pial surface (boundary between grey matter and cerebrospinal fluid) and the white surface (boundary between grey matter and white matter), as detailed in Figure 3.1.

3.2 Functional Magnetic Resonance Imaging

Ideally, in cognitive neuroscience we would like to observe neuronal activity directly and relate it to ongoing cognitive processes as for example probed in closely controlled experimental studies. However, this is not yet possible in a non-invasive fashion today. Instead, we use a proxy to neuronal activity, the so-called Blood Oxygen Level Dependent (BOLD) contrast (Ogawa, Lee, Kay, & Tank, 1990). This contrast relies on the fact that the deoxyhemoglobin concentration in the blood changes when more oxygen is being consumed, for example because the participant is currently doing a decision-making task and active brain areas need to be supplied with more oxygen. This physiological change in deoxyhemoglobin concentration changes the magnetic properties of water molecules which is then measurable by MRI. The concept of BOLD-based functional MRI (fMRI) relies on this indirect measure to assess neural activity in the form of local field potentials (Logothetis, Pauls, Augath, Trinath, & Oeltermann, 2001). We make use of fMRI and the BOLD contrast in particular in **Study I** using Statistical Parametric Mapping (SPM, Wellcome Trust Centre for Neuroimaging, UCL, London, United Kingdom) as the analysis tool.

3.3 Reinforcement Learning Modelling

We used a standard RL model in **Study I** to investigate participant's choices and their learning progress. Because we assumed that participants learn differently from positive and negative feedback, we estimated two learning rates, α_{pos} and α_{neg} , respectively. We initialized all weights at 0, the two learning rates at 0.5 (with $0 \leq \alpha \leq 1$ and a beta distributed prior) and the inverse temperature parameter β at 1 (with $\beta \geq 0$ and a normal distributed prior with mean = 0 and standard deviation = $\sqrt{10}$).

On a trial-by-trial basis, stimulus values were estimated via the following formula:

$$Q_{t+1} = Q_t + \alpha_{(pos/neg)} * \delta_t$$

with the prediction errors δ_t calculated as:

$$\delta_t = r_t - Q_t$$

Action probabilities (here for symbol A in pair AB) were estimated using the softmax equation:

$$P(A)_t = \frac{1}{1 + e^{-\beta(Q(A)_t - Q(B)_t)}}$$

Further, we calculated a trial-by-trial confidence measure by putting the action tendency estimated in the softmax equation in relation to 0.5:

$$Conf(AB)_t = |0.5 - P(A)_t|$$

3.4 Hierarchical Bayesian Estimation of the Drift Diffusion Model (HDDM)

To model latent decision processes of PES in **Studies I** and **II** we used a toolbox which allows the estimation of the drift diffusion model in a hierarchical Bayesian fashion (HDDM, Wiecki, Sofer, & Frank, 2013). In contrast to traditional ways of estimating diffusion models, the hierarchical Bayesian analysis assumes that parameters for individuals can be drawn from the group distribution and that an uncertainty around the parameters can also be estimated. A major benefit of this type of analysis is that far fewer trials per condition are needed for each individual participant to get reasonable parameter estimates as the individual parameter estimates are constrained by the group estimates (Ratcliff & Cholders, 2015; Wiecki et al., 2013). In our studies, we mainly focused on the three main parameters of the DDM: The decision threshold a , the rate of evidence accumulation v , and the non-decision time T_{er} . The reaction time slowing after an error can be described by either an increase in decision threshold, a decrease in evidence accumulation or an increase in the non-decision time, see Figure 3.2. Teasing these contributions apart was particularly our goal in **Study III**.

3.5 Post-error slowing

Traditionally, post-error slowing has been calculated by comparing average RT on post-error trials with average RT on post-correct trials. However, this method neglects the fact that errors might also cluster in certain parts of the experiment, e.g., at the end when general attention is low because participants are getting tired. It is therefore advisable to take this into account and subtract post-error RTs from associated pre-error RTs (so-called Δ RT) instead so that global fluctuations over

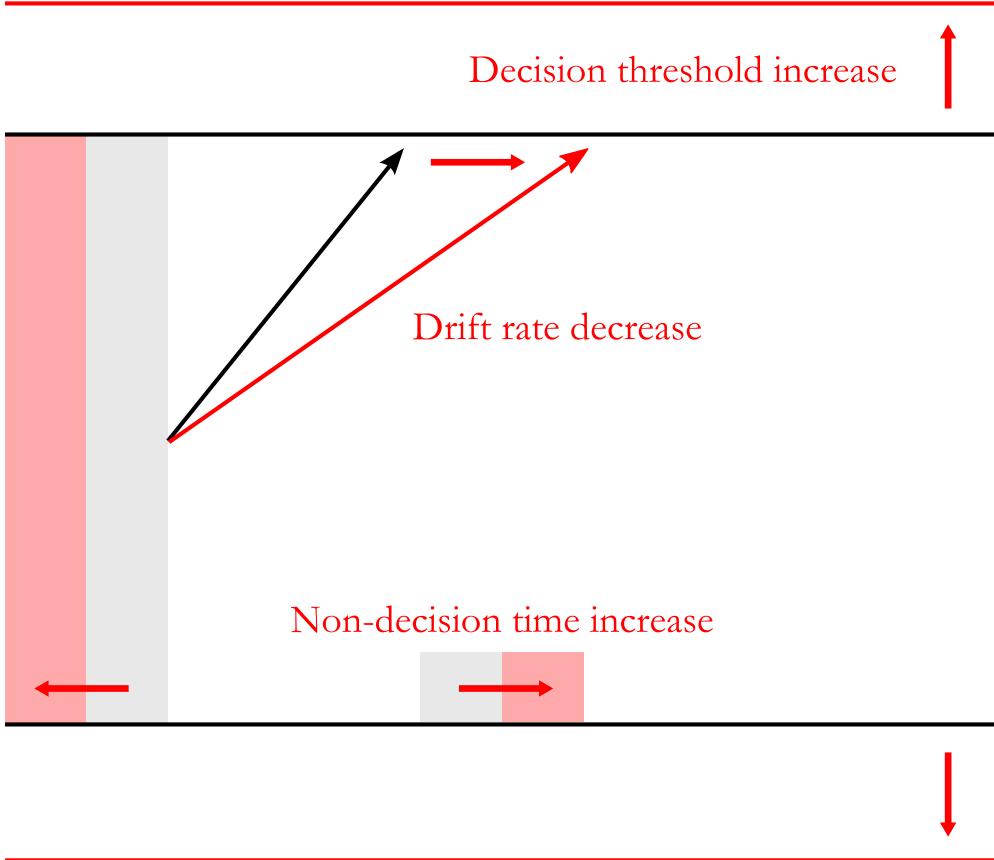


Figure 3.2: Purported decision process changes underlying post-error slowing. All three changes portrayed lead to a decrease in RT. However, changes in the respective parameters of the drift diffusion model map uniquely to changes in accuracy. While decision threshold increases predict corresponding increases in accuracy, decreases in evidence accumulation will lead to lower accuracy, and changes in non-decision time do not affect accuracy.

the course of the experiment are being taken into account (Dutilh et al., 2012a). In our studies, we also control for trial type differences in RT by comparing trials of the same trial condition. For **Studies I** and **II**, we calculate PES in accordance with previous research on post-error slowing in a reinforcement learning task design (Cavanagh et al., 2010; Frank et al., 2007) by subtracting post-error RT for the next same pair from RT on the error trial. For **Study III**, we calculate PES by subtracting post-error RT from pre-error RT if they are of the same trial type (neutral/angry/happy).

3.6 Participants

For **Studies I** and **II** we recruited a total of 48 healthy participants who gave written informed consent before participating in the study. In **Study III**, 6,047 participants took part in the study which was presented at an art exhibition at Nationalmuseum in Stockholm, Sweden. Detailed information about the aims of the research was given to participants both on the TV screen for the experiment and on text panels of the installation and consent was implied by participants voluntarily initiating the task.

Chapter 4

Results

4.1 Study I: Post-error slowing is associated with learning performance and functional activity in cognitive control and visual regions

4.1.1 Relation of learning phase measures to the testing phase

In a multiple regression model, we found that both learning phase accuracy on the main symbol pair (AB) and PES during the learning phase were positively associated with testing phase performance (Correlation between PES during learning phase and test score in Figure 4.1A), while the number of times participants decided to stay with the same decision after negative feedback or switch to the other option was not directly related to test phase performance.

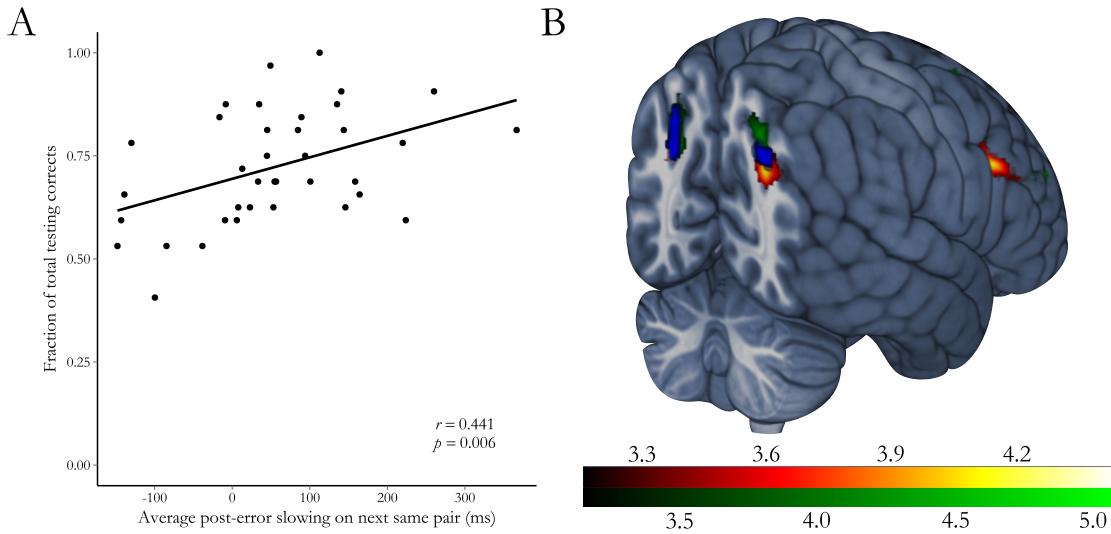


Figure 4.1: Main results of Study I. (A) Positive correlation between memory-based post-error slowing and test phase performance across participants. (B) fMRI activity of brain areas while receiving negative feedback, related to amount of slowing on the next same pair trial (red-yellow) and to absolute prediction error (green), as well as the conjunction between both (blue). Figures from Schiffler et al., 2016, reprinted with permission from MIT Press.

We did not find a correlation of PES with overall accuracy of any of the symbol pairs during the learning phase. Testing phase scores also demonstrated that participants performed better in the test phase on new symbol combinations which were easier. For example, the choice between symbol A which was reinforced at an 80% probability during the learning phase and symbol D, which was reinforced at 30% should be easier than the choice between A and C (80%/70%), see Figure 4.2.

4.1.2 Feedback-congruent staying/shifting

As expected, participants on average repeated decisions more often when they were reinforced by positive feedback compared to negative feedback. A general working memory component as indicated by feedback congruent behaviour in the beginning of the task (Frank et al., 2007) was not significantly related to PES.

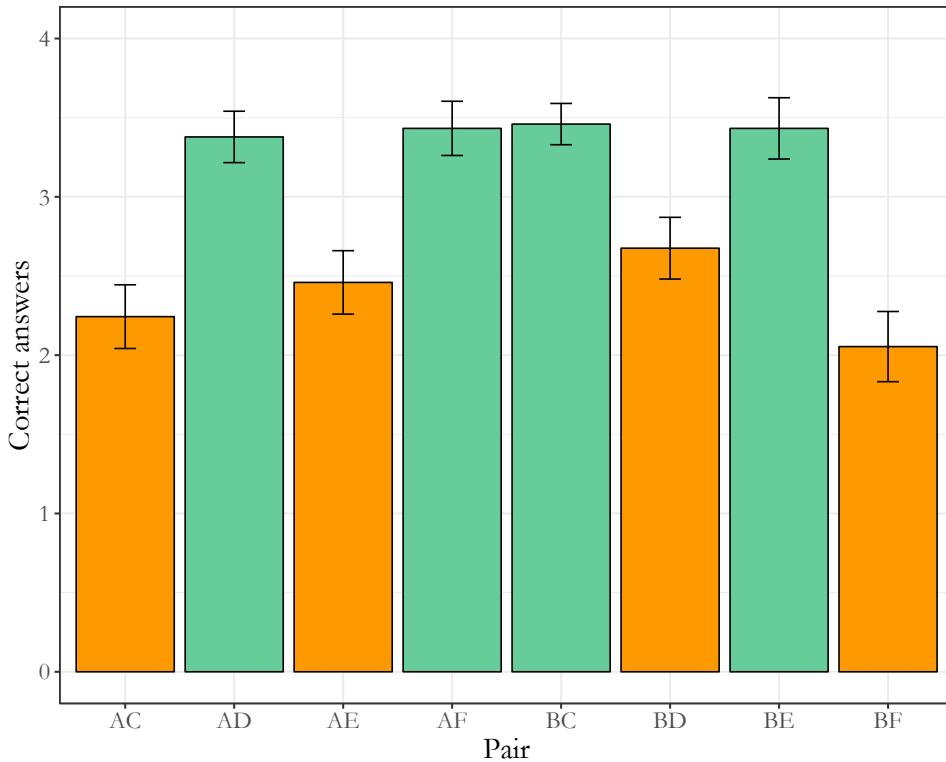


Figure 4.2: Test scores divided by symbol pair combinations. During the test phase of the task, the symbols A and B from the learning phase are now tested separately against all other symbols learned. Green colour represents easier symbol combinations (AD: 80%/30% probability of positive reinforcement, AF: 80%/40%, BC: 20%/70%, BE: 20%/60%) and orange represents more difficult combinations (AC: 80%/70%, AE: 80%/60%, BD: 20%/30%, BF: 20%/40%). Error bars reflect standard error of the mean.

4.1.3 fMRI activity associated with PES

On the event of receiving negative feedback, activity in right inferior frontal cortex and bilateral occipital cortex tracked the response speed adjustment on the following relevant trial (Figure 4.1B).

4.1.4 Reinforcement learning model measures

4.1.4.1 Prediction errors and their neural correlates

Prediction errors estimated from our reinforcement learning model were associated with post-error slowing. More unexpected negative feedback lead to an increase in slowing both on the direct next trial and on the next relevant (same pair) trial while more unexpected positive feedback was followed by a speed increase. Negative prediction errors correlated with brain activity in left striatum as assessed by an *a priori* ROI analysis and absolute (i.e., unsigned) prediction errors over all feedback evoked activity in similar brain regions as the main PES analysis.

4.1.4.2 Learning rate

We had estimated two separate learning rates for positive and negative feedback in accordance with previous research (Kahnt et al., 2009; Van Den Bos et al., 2012) to investigate whether participants who showed stronger reactivity to negative feedback (e.g., increased slowing or switching to the other symbol) showed a learning pattern which focuses on recent feedback in contrast to the whole history (i.e., a high learning rate). We found this pattern for switch behaviour, but not for post-error slowing. This means that the model estimated a higher negative learning rate for participants who switched their choice to the other symbol following negative feedback.

4.1.4.3 Confidence

Confidence measures as extracted by our RL model also showed a negative relation to PES (i.e., lower confidence lead to an increase in ΔRT).

4.1.4.4 Model validation

We simulated model predictions by taking the final model parameters ($\alpha_{pos,neg}$ and β) and averaged over 10,000 repetitions of simulated behaviour for each participant. The model was able to reproduce the learning curves present in the acquired data as well as the differentiation in accuracy between the three different symbol pairs at the end of the training phase.

4.2 Study II: Adaptive increases in response caution after errors are related to cortical thickness in cognitive control regions

4.2.1 Drift diffusion correlates of PES in a reinforcement learning design

Post-error trials, compared to post-correct trials, were defined by an increase in decision threshold but there was no difference in the rate of evidence accumulation (Figure 4.3A,B).

This indicates that error trials evoke increased response caution in this task design but not necessarily decreases in evidence accumulation. Against our expectations, an interaction between the distance to the next same pair trial and the rate of evidence accumulation was not supported by the data. These findings suggest that memory-based PES is an adaptive cognitive control process because the RT increases in this experiment contributed to accuracy as shown by the increase in response caution.

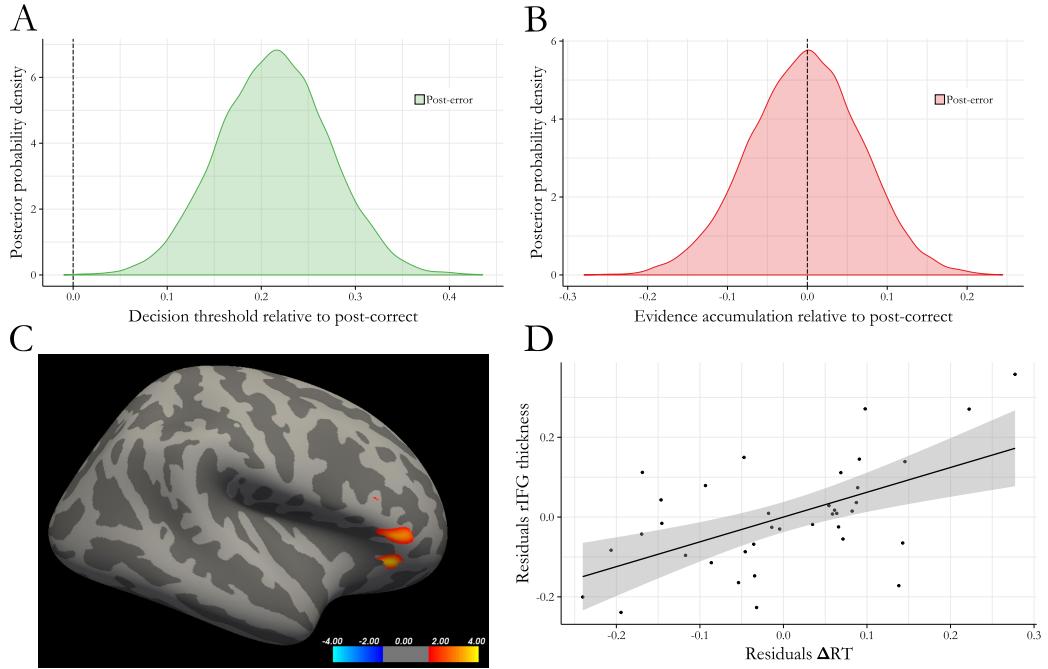


Figure 4.3: Main results of Study II. (A) Posterior probability density of decision threshold parameter regression estimate for post-error trials compared to post-correct trials. (B) Posterior probability density of evidence accumulation parameter regression estimate for post-error trials compared to post-correct trials. (C) Vertex based analysis showing clusters in which rIFG thickness is positively associated with ΔRT . (D) Correlation between average ΔRT and average rIFG cortical thickness when partialing out the effects of age, sex and prime. Shaded area indicates 95% confidence interval of the regression line.

4.2.2 PES is related to cortical thickness in rIFC

We found that overall cortical thickness of rIFC correlated with PES, both if the same pair was the direct next trial after the error and if there was at least one other symbol pair in between. *Post-hoc* correlations and a follow-up vertex wise analysis both demonstrated that the strongest association of PES was to the anterior part of rIFC (Figure 4.3C,D), in pars orbitalis (particularly for longer distances) and pars triangularis (especially for immediate next trials).

4.2.3 Response caution increases after errors directly relate to cortical thickness in rIFC

Using participants' parameters of decision threshold and evidence accumulation adaptation on post-error trials compared to post-correct trials, we showed that anatomical variability in rIFC, particularly in pars orbitalis, related to decision threshold adaptations on the trial after the error.

4.3 Study III: Response adaptations to errors are multi-component processes and change dynamically over several trials after the error

4.3.1 Dynamics of latent decision processes after errors

In this study, we have found that reactions are marked by sustained increases in response caution over several trials and transient decreases in evidence accumulation, most prominent on the direct next trial after the error (Figure 4.4A).

Further, we also found that non-decision time was reduced for several trials post-error. This indicates that in our experiment, multiple decision processes were affected simultaneously by the error and give rise to the particular pattern of post-error slowing.

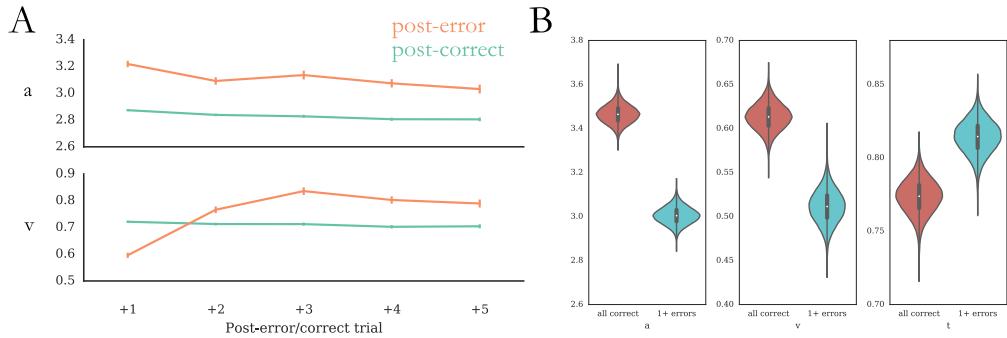


Figure 4.4: Main results of Study III. (A) Decision process components contributing to post-error slowing over several trials after an error. While decision threshold a showed a sustained increase even several trials after an error (orange) compared to post-correct trials (green), an initial post-error dip in evidence accumulation v increased over the following trials. Non-decision time T_{er} (not shown) displayed a sustained decrease for post-error trials relative to post-correct trials. (B) Relation of decision process parameters on the first trial after an error to accuracy on the following five trials. Higher decision thresholds a and evidence accumulation v as well as lower non-decision time T_{er} immediately post-error were associated with no mistakes on the following five trials compared to at least one mistake. Figures from Schiffler et al., 2017, adapted under CC BY 4.0 license.

4.3.2 How do downstream accuracy increases rely on post-error adaptation?

Post-error trials which were followed by five correct trials were marked by higher decision thresholds, higher rates of evidence accumulation and lower non-decision times (Figure 4.4B).

4.3.3 Effects of error properties on decision making

We found differences in post-error decision processes both by emotional valence and difficulty of stimuli. These variations had specific effects on the rate of evidence accumulation. It was reduced both for generally more difficult trials and following angry error trials compared to happy error trials.

Chapter 5

Discussion

Can appropriate response adaptation to negative feedback support learning? Our studies presented in this thesis suggest that this is indeed generally the case, but that there are also potentially detrimental components which need to be recognized.

5.1 Behavioural results in Studies I - III

In **Studies I** and **II** we have presented evidence that PES in a RL context can benefit learning outcome. **Study I** demonstrated that there are long-term learning outcome benefits if participants adapt their response speed to the negative feedback received. Interestingly, the effect that we found for PES on testing phase performance was similar in magnitude to the effect of overall learning phase performance in the main symbol pair AB on the test score. This is somewhat surprising as being able to clearly differentiate between the value of these two symbols as indicated by accuracy on pair AB during the learning phase should be a strong predictor for performance during the test phase in which symbols A and B are separately pitted against all other symbols.

These findings extend the literature on beneficial and adverse effects of PES on learning by showing that error monitoring processes after negative feedback can be an indicator for later learning outcome as stipulated by Ridderinkhof and colleagues (2004). This suggests that negative feedback is encoded at the time of receiving the feedback to drive future response adaptations. This finding emphasizes encoding of feedback as a crucial phase for subsequent changes in behaviour, which may have implications for structuring learning exercises, e.g., in the classroom.

In **Study II**, we showed that PES also has a positive effect during the initial learning phase of the task. The increase in RT was associated with an increase in the decision threshold, which signals that participants made more cautious (and thus on average more correct) decisions after errors. This finding is not very surprising, given that post-error adaptations in general have previously been associated with increases in response caution (Dutilh et al., 2012b) and that in a different reinforcement learning task, decision threshold has been linked to decision conflict (Frank et al., 2015). The result reiterates that the additional time that is often being taken after errors is beneficial to learning performance.

Study III provided evidence for both functional and potentially maladaptive decision components in the process of post-error slowing. While response caution persisted at elevated levels even several trials after the error, reductions in evidence accumulation were mainly present for the first post-error trial, suggesting only a transient disruption of decision making by the error.

Reduced evidence accumulation and increased decision boundaries have also been found in a motion direction discrimination task (Purcell & Kiani, 2016), although that study did not find any differences in non-decision time. However, in our case, the non-decision time was reduced following errors, i.e., this part of the decision process did not contribute to post-error slowing.

Our findings align with a recent conceptual proposal, which suggests both

co-occurring increases in response threshold and decreases in sensitivity to sensory information which together lead to initial decreases in accuracy and eventual recovery over future trials (Ullsperger & Danielmeier, 2016).

Interestingly, these results as predicted by the aforementioned theory stand in contrast to Laming's (1979) original findings that RT recovers faster after an error than task accuracy, possibly due to differences in the task being used or in the way that post-error slowing was calculated.

5.2 Neuroimaging results in Studies I and II

In **Studies I and II**, we show that PES was associated with both function and anatomy of the rIFG, an important cognitive control region in the brain. Specifically, anatomical variability in cortical thickness in this area reflected participants' decision threshold adaptations after errors, suggesting a possibility that there are markers of propensity to error adaptation in the human brain which can be investigated in further research. Functional activity in rIFG when receiving the error feedback was positively associated with response time slowing on the next relevant trial.

In similar previous analyses, activity in dorsolateral PFC (Kerns et al., 2004) and pMFC (Danielmeier et al., 2011) on the error trial predicted RT slowing on the post-error trial. Hester and colleagues also demonstrated that error activity in pMFC predicted accuracy on the next same target stimulus, even if that trial was several trials ahead in the future (Hester, Madeley, Murphy, & Mattingley, 2009). Further, activity in bilateral IFG has previously been associated with successful instrumental learning (Guitart-Masip et al., 2012).

It is unlikely that there exists only one single region in PFC which determines

response adaptations like PES, but that instead several areas cooperate to implement the response speed adjustments in cooperation with primary motor areas, the pre-SMA and subcortical regions like the STN (Siegert et al., 2014). Further, these areas likely interact with regions which have been consistently shown to be involved in error monitoring like pMFC and ACC. Our contribution in the studies presented here is that lateral PFC, and more specifically rIFG, is directly involved in signalling the need for a future more cautious response after receiving the feedback, not only for suppressing the motor output on the post-error trial. Future studies which focus on functional or anatomical network approaches (as e.g., in Rae et al., 2015) will be able to provide a more comprehensive answer to these questions.

5.3 Limitations and future directions

Particularly for our drift diffusion modelling we have employed a method which enables accurate parameter estimation even when trial amounts are low. However, for the particular conditions we wanted to investigate, we still encountered problems in convergence and had to make compromises on the model structure (i.e., model the data on group level in **Study III** and reduce model complexity in **Study II**) to assure convergence of the models. Even though we had a lot of participants in **Study III** and the direction of almost all parameter estimates when including individual level modelling was the same as for the group models, the trial amount per participant should still be increased to also get stable individual parameter estimates.

Further, it is important to consider that the PES effects we have observed here differ between the presented studies due to the intricacies of the two experimental designs. Particularly, this concerns the memory-based aspect of PES in **Study I** and **II**

compared to a more conventional visual search task in **Study III**. While adaptation of response speed in line with previous relevant feedback is reasonable and has previously been found in other studies employing the same type of probabilistic reinforcement learning task (Cavanagh et al., 2010; Frank et al., 2007), this might constitute a different aspect of PES than what has historically been classified as PES. For example, one recurring finding in PES research (Danielmeier & Ullsperger, 2011; see e.g., Jentzsch & Dudschg, 2009) is that the slowing is greater when the response to stimulus interval is smaller (i.e., the closer the time from feedback to onset of next stimulus). In our first two studies, a considerable amount of time could pass between a particular feedback and the next relevant trial (on average 20 seconds between feedback and next same pair in comparison to traditional studies of PES in which the response to stimulus interval is usually below one second). As such, the received negative feedback needs to be stored in some way until the participant recognizes the same pair for the next time. What we investigate in the first two studies could thus involve a more cognitive aspect of post-error slowing than has been conventionally examined.

In **Study I** and **II**, the PES effect was also comparatively smaller to PES in **Study III**. This might be because there were few error trials in the latter study which bias the analysis of post-error reactions towards initial encounters of errors. Conceivably, these initial errors might provoke a stronger post-error reaction, but this hypothesis would need to be tested in future studies.

Another interesting aspect for future research concerns the question how errors in deterministic decision-making tasks differ from probabilistic negative feedback often given in reinforcement learning task designs with regard to involvement of brain areas associated with cognitive control. One possibility is that while the former feedback evokes a similar response independent of the current stage of the task, the latter might lead to decreasing engagement of relevant cognitive control structures when the value of a particular stimulus is already certain (e.g., in later

phases of a task).

Finally, I believe that both the fields of cognitive control and reinforcement learning will benefit immensely from research on dynamic functional connectivity (R. M. Hutchison et al., 2013; Thompson, Brantefors, & Fransson, 2017). The promise of focusing on the dynamics of brain activity lies in the potential to elucidate the brain networks contributing to processes such as post-error slowing on a moment-to-moment basis and how these networks interact with other brain areas to promote learning.

For the studies presented here, this would mean that the findings about rIFG and potential feature processing regions in occipital cortex could be embedded in a larger context of how they work together with other nodes like medial PFC, STN and pre-SMA in a network which regulates decision threshold adaptations (Cavanagh, Sanguinetti, Allen, Sherman, & Frank, 2014; Cavanagh et al., 2011; Frank et al., 2015; Herz, Zavala, Bogacz, & Brown, 2016; Rae et al., 2015) and with areas in the brain encoding the value or deviance from expected value of relevant stimuli such as striatal areas like the nucleus accumbens (Niv et al., 2012).

A concrete example of how the interaction of those networks could be probed is by combining experimental paradigms which have usually been used to evoke specific brain activity in relation to cognitive control and reinforcement learning. For example, it could be possible to first “load” certain stimuli (e.g., faces and houses) with high and low value by reinforcing them and then in a second phase use these stimuli in e.g., a Go/No-Go task to determine how the previous reinforcement affects cognitive control performance (see e.g., Freeman, Razhas, & Aron, 2014 for an example of combining a conditioning task with a Go/No-Go task). With proper orthogonalization of stimuli (e.g., using dimensions of face characteristics and type of house in the Go/No-Go task), specific hypotheses about the involved brain networks could then be tested. Of particular interest could

be an interaction of cognitive control network areas with both reward related networks and stimulus-specific association parts of the brain (Danielmeier et al., 2011; Schiffer, Muller, Yeung, & Waszak, 2014). This would provide a more comprehensive picture of the role of value in cognitive control. It could potentially even be used to inform research on clinical disorders such as addiction by asking why prepotent action tendencies associated with reward can sometimes not easily be controlled, verifying earlier findings of involvement of rIFC in inhibiting craving (Tabibnia et al., 2011; Volkow et al., 2010).

Returning to Foster's original vision to decompose reaction times into decision components, we now have the computational tools available to do so on increasingly larger amounts of data. For example in clinical populations, making use of that computational power to investigate differences in decision processes promises substantial information gain compared to earlier, more coarse approaches and can aid in bringing forth the emerging field of computational psychiatry (Huys, Maia, & Frank, 2016; Maia & Frank, 2011; Montague, Dolan, Friston, & Dayan, 2012).

Acknowledgments

First of all, I would like to say an enormous thank you to my supervisor **Sara Bengtsson** for taking me on as your PhD student. You have always had an open ear and door for me throughout this whole PhD and I am very grateful for that!

Thank you also to my two co-supervisors, **Martin Ingvar** and **Peter Fransson**. **Martin**, I sincerely respect your scientific rigour and enthusiasm. **Peter**, I was very thankful for being able to pick your brain on various methodological questions related to neuroimaging.

To **Granville**, **Lieke**, **Nina**, **Pontus**, and **William**, my closest support group during the PhD: Our collaborations, journal clubs and activities outside of work made my sceptical mind curious about science again when it was needed the most. Thank you so much for everything!

To **Pär** and **Alva**, you have given me a better welcome in Stockholm than I could have wished for and were a constant source of encouragement, particularly in the difficult phases at the beginning of the PhD.

Anaïs, thank you for being a great roommate for almost the entire PhD time and **Fanny**, thank you for the tennis matches and the culinary support of our corridor!

To **Mimmi**, with your positive spirit you always managed to turn big administrative hassles into minor inconveniences.

A special thanks to **Rita** for great statistical aid (often with short deadlines) and to **Karin** for long discussions about the newest artful TV series and the support towards the end of the PhD.

To all the other people in the old and new corridor who have made life at KI so much more enjoyable: **Frida, Sofia, Christoph, Eleni, Jonathan, Aisha, Jeanette, Sandra, Gustav, Nathalie, Isabel, Eva, Predrag, Angelica, Mikkel, Lau, Cassia, Daniel, Annelie, Orestis, Miriam, Diana, Emilia, and Patrik.**

Thank you also to the other members in the various journal clubs from whom I have learnt a lot, particularly **Philip** and **Ida**.

To my Stockholm friends **Benji, Gustaf & Elsa, Anirudh**, and my family away from home: **Margareta & Per-Olov**.

Thank you to my friends back in Germany who have been with me for a long time: **Markus S., Markus H., Thorsten, and Andi**.

My dear **Lu**, you have changed my life for a much better course. I am so happy to have you by my side and look forward to our future adventures! 谢谢你!

Zum Schluss vielen lieben Dank an meine Eltern **Helga** und **Helmut** und meinen Bruder **Max** für eure Unterstützung über die Jahre!

References

- Aron, A. R., Behrens, T. E., Smith, S., Frank, M. J., & Poldrack, R. A. (2007). Triangulating a Cognitive Control Network Using Diffusion-Weighted Magnetic Resonance Imaging (MRI) and Functional MRI. *Journal of Neuroscience*, 27(14), 3743–3752. <https://doi.org/10.1523/JNEUROSCI.0519-07.2007>
- Aron, A. R., Fletcher, P. C., Bullmore, E. T., Sahakian, B. J., & Robbins, T. W. (2003). Stop-signal inhibition disrupted by damage to right inferior frontal gyrus in humans. *Nature Neuroscience*, 6(2), 115–116. <https://doi.org/10.1038/nn1003>
- Aron, A. R., Robbins, T. W., & Poldrack, R. A. (2014). Inhibition and the right inferior frontal cortex: One decade on. *Trends in Cognitive Sciences*, 18(4), 177–185. <https://doi.org/10.1016/j.tics.2013.12.003>
- Boekel, W., Forstmann, B. U., & Keuken, M. C. (2017). A test-retest reliability analysis of diffusion measures of white matter tracts relevant for cognitive control. *Psychophysiology*, 54(1), 24–33. <https://doi.org/10.1111/psyp.12769>
- Bogacz, R., Wagenmakers, E.-J., Forstmann, B. U., & Nieuwenhuis, S. (2010). The neural basis of the speed-accuracy tradeoff. *Trends in Neurosciences*, 33(1), 10–16. <https://doi.org/10.1016/j.tins.2009.09.002>
- Botvinick, M. M., & Weinstein, A. (2014). Model-based hierarchical reinforcement learning and human action control. *Philosophical Transactions of the Royal Society*

B: *Biological Sciences*, 369(1655), 20130480. <https://doi.org/10.1098/rstb.2013.0480>

Botvinick, M. M., Braver, T. S., Barch, D. M., Carter, C. S., & Cohen, J. D. (2001). Conflict monitoring and cognitive control. *Psychological Review*, 108(3), 624–652. <https://doi.org/10.1037/0033-295X.108.3.624>

Braver, T. S. (2012). The variable nature of cognitive control: A dual mechanisms framework. *Trends in Cognitive Sciences*, 16(2), 106–113. <https://doi.org/10.1016/j.tics.2011.12.010>

Cavanagh, J. F., Frank, M. J., Klein, T. J., & Allen, J. J. B. (2010). Frontal theta links prediction errors to behavioral adaptation in reinforcement learning. *NeuroImage*, 49(4), 3198–209. <https://doi.org/10.1016/j.neuroimage.2009.11.080>

Cavanagh, J. F., Sanguinetti, J. L., Allen, J. J. B., Sherman, S. J., & Frank, M. J. (2014). The Subthalamic Nucleus Contributes to Post-error Slowing. *Journal of Cognitive Neuroscience*, 26(11), 2637–2644. https://doi.org/10.1162/jocn_a_00659

Cavanagh, J. F., Wiecki, T. V., Cohen, M. X., Figueroa, C. M., Samanta, J., Sherman, S. J., & Frank, M. J. (2011). Subthalamic nucleus stimulation reverses mediofrontal influence over decision threshold. *Nature Neuroscience*, 14(11), 1462–7. <https://doi.org/10.1038/nn.2925>

Chambers, C. D., Bellgrove, M. A., Stokes, M. G., Henderson, T. R., Garavan, H., Robertson, I. H., ... Mattingley, J. B. (2006). Executive “Brake Failure” following Deactivation of Human Frontal Lobe. *Journal of Cognitive Neuroscience*, 18(3), 444–455. <https://doi.org/10.1162/jocn.2006.18.3.444>

Chase, H. W., Kumar, P., Eickhoff, S. B., & Dombrovski, A. Y. (2015). Reinforcement learning models and their neural correlates: An activation

likelihood estimation meta-analysis. *Cognitive, Affective, & Behavioral Neuroscience*, 15(2), 435–59. <https://doi.org/10.3758/s13415-015-0338-7>

Cohen, J. D., McClure, S. M., & Yu, A. J. (2007). Should I stay or should I go? How the human brain manages the trade-off between exploitation and exploration. *Philosophical Transactions of the Royal Society of London. Series B, Biological Sciences*, 362(1481), 933–42. <https://doi.org/10.1098/rstb.2007.2098>

Danielmeier, C., & Ullsperger, M. (2011). Post-error adjustments. *Frontiers in Psychology*, 2. <https://doi.org/10.3389/fpsyg.2011.00233>

Danielmeier, C., Eichele, T., Forstmann, B. U., Tittgemeyer, M., & Ullsperger, M. (2011). Posterior Medial Frontal Cortex Activity Predicts Post-Error Adaptations in Task-Related Visual and Motor Areas. *Journal of Neuroscience*, 31(5), 1780–1789. <https://doi.org/10.1523/JNEUROSCI.4299-10.2011>

Daw, N. D. (2011). Trial-by-trial data analysis using computational models. *Decision Making, Affect, and Learning: Attention and Performance XXIII*, 23, 3–38. <https://doi.org/10.1093/acprof:oso/9780199600434.003.0001>

Dickerson, B. C., & Wolk, D. A. (2012). MRI cortical thickness biomarker predicts AD-like CSF and cognitive decline in normal adults. *Neurology*, 78(2), 84–90. <https://doi.org/10.1212/WNL.0b013e31823efc6c>

Dickerson, B. C., Bakkour, A., Salat, D. H., Feczkó, E., Pacheco, J., Greve, D. N., ... Buckner, R. L. (2009). The cortical signature of Alzheimer's disease: Regionally specific cortical thinning relates to symptom severity in very mild to mild AD dementia and is detectable in asymptomatic amyloid-positive individuals. *Cerebral Cortex*, 19(3), 497–510. <https://doi.org/10.1093/cercor/bhn113>

Doll, B. B., Simon, D. A., & Daw, N. D. (2012). The ubiquity of model-based reinforcement learning. *Current Opinion in Neurobiology*, 22(6), 1075–1081. <https://doi.org/10.1016/j.conb.2012.05.003>

//doi.org/10.1016/j.conb.2012.08.003

Dutilh, G., Van Ravenzwaaij, D., Nieuwenhuis, S., Van der Maas, H. L. J., Forstmann, B. U., & Wagenmakers, E.-J. (2012a). How to measure post-error slowing: A confound and a simple solution. *Journal of Mathematical Psychology*, 56(3), 208–216. <https://doi.org/10.1016/j.jmp.2012.04.001>

Dutilh, G., Vandekerckhove, J., Forstmann, B. U., Keuleers, E., Brysbaert, M., & Wagenmakers, E.-J. (2012b). Testing theories of post-error slowing. *Attention, Perception, & Psychophysics*, 74(2), 454–465. <https://doi.org/10.3758/s13414-011-0243-2>

Economides, M., Guitart-Masip, M., Kurth-Nelson, Z., & Dolan, R. J. (2014). Anterior Cingulate Cortex Instigates Adaptive Switches in Choice by Integrating Immediate and Delayed Components of Value in Ventromedial Prefrontal Cortex. *Journal of Neuroscience*, 34(9), 3340–3349. <https://doi.org/10.1523/JNEUROSCI.4313-13.2014>

Fischl, B., & Dale, A. M. (2000). Measuring the thickness of the human cerebral cortex from magnetic resonance images. *Proceedings of the National Academy of Sciences*, 97(20), 11050–11055. <https://doi.org/10.1073/pnas.200033797>

Fitts, P. M. (1966). Cognitive aspects of information processing. 3. Set for speed versus accuracy. *Journal of Experimental Psychology*, 71(6), 849–857. <https://doi.org/10.1037/h0023232>

FitzGerald, T. H. B., Friston, K. J., & Dolan, R. J. (2012). Action-Specific Value Signals in Reward-Related Regions of the Human Brain. *Journal of Neuroscience*, 32(46), 16417–16423. <https://doi.org/10.1523/JNEUROSCI.3254-12.2012>

Fjell, A. M., Westlye, L. T., Amlie, I. K., & Walhovd, K. B. (2012). A multi-modal investigation of behavioral adjustment: Post-error slowing is associated with white

matter characteristics. *NeuroImage*, 61(1), 195–205. <https://doi.org/10.1016/j.neuroimage.2012.03.007>

Forstmann, B. U., Ratcliff, R., & Wagenmakers, E.-J. (2016). Sequential Sampling Models in Cognitive Neuroscience: Advantages, Applications, and Extensions. *Annual Review of Psychology*, 67(1), 641–666. <https://doi.org/10.1146/annurev-psych-122414-033645>

Foster, M. (1870). The Velocity of Thought. *Nature*, 2, 2–4.

Frank, G. K. W., Reynolds, J. R., Shott, M. E., & O'Reilly, R. C. (2011). Altered temporal difference learning in bulimia nervosa. *Biological Psychiatry*, 70(8), 728–735. <https://doi.org/10.1016/j.biopsych.2011.05.011>

Frank, M. J., Gagne, C., Nyhus, E., Masters, S., Wiecki, T. V., Cavanagh, J. F., & Badre, D. (2015). fMRI and EEG Predictors of Dynamic Decision Parameters during Human Reinforcement Learning. *The Journal of Neuroscience*, 35(2), 485–494. <https://doi.org/10.1523/JNEUROSCI.2036-14.2015>

Frank, M. J., Moustafa, A. A., Haughey, H., Curran, T., & Hutchison, K. (2007). Genetic triple dissociation reveals multiple roles for dopamine in reinforcement learning. *Proceedings of the National Academy of Sciences of the United States of America*, 104(41), 16311–6. <https://doi.org/10.1073/pnas.0706111104>

Frank, M. J., Seeberger, L. C., & O'Reilly, R. C. (2004). By carrot or by stick: cognitive reinforcement learning in parkinsonism. *Science*, 306(5703), 1940–1943. <https://doi.org/10.1126/science.1102941>

Freeman, S. M., Razhas, I., & Aron, A. R. (2014). Top-down response suppression mitigates action tendencies triggered by a motivating stimulus. *Current Biology*, 24(2), 212–216. <https://doi.org/10.1016/j.cub.2013.12.019>

Friston, K. J., & Kiebel, S. (2009). Predictive coding under the free-energy

principle. *Philosophical Transactions of the Royal Society of London. Series B, Biological Sciences*, 364(1521), 1211–21. <https://doi.org/10.1098/rstb.2008.0300>

Garrison, J., Erdeniz, B., & Done, J. (2013). Prediction error in reinforcement learning: A meta-analysis of neuroimaging studies. *Neuroscience and Biobehavioral Reviews*, 37(7), 1297–1310. <https://doi.org/10.1016/j.neubiorev.2013.03.023>

Gläscher, J., Adolphs, R., Damasio, H., Bechara, A., Rudrauf, D., Calamia, M., ... Tranel, D. (2012). Lesion mapping of cognitive control and value-based decision making in the prefrontal cortex. *Proceedings of the National Academy of Sciences of the United States of America*, 109(36), 14681–6. <https://doi.org/10.1073/pnas.1206608109>

Gläscher, J., Daw, N. D., Dayan, P., & O'Doherty, J. P. (2010). States versus rewards: Dissociable neural prediction error signals underlying model-based and model-free reinforcement learning. *Neuron*, 66(4), 585–595. <https://doi.org/10.1016/j.neuron.2010.04.016>

Guitart-Masip, M., Huys, Q. J. M., Fuentemilla, L., Dayan, P., Duzel, E., & Dolan, R. J. (2012). Go and no-go learning in reward and punishment: Interactions between affect and effect. *NeuroImage*, 62(1), 154–166. <https://doi.org/10.1016/j.neuroimage.2012.04.024>

Hajcak, G., McDonald, N., & Simons, R. F. (2003). To err is autonomic: Error-related brain potentials, ANS activity, and post-error compensatory behavior. *Psychophysiology*, 40(6), 895–903. <https://doi.org/10.1111/1469-8986.00107>

Heitz, R. P. (2014). The speed-accuracy tradeoff: History, physiology, methodology, and behavior. *Frontiers in Neuroscience*, 8. <https://doi.org/10.3389/fnins.2014.00150>

Helmholtz, H. von. (1850). Messungen über den zeitlichen Verlauf der Zuckung

animalischer Muskeln und die Fortpflanzungsgeschwindigkeit der Reizung in den Nerven. *Archiv Für Anatomie, Physiologie Und Wissenschaftliche Medicin*, 276–364.

Herz, D. M., Zavala, B. A., Bogacz, R., & Brown, P. (2016). Neural Correlates of Decision Thresholds in the Human Subthalamic Nucleus. *Current Biology*, 26(7), 916–920. <https://doi.org/10.1016/j.cub.2016.01.051>

Hester, R., Barre, N., Mattingley, J. B., Foxe, J. J., & Garavan, H. (2007). Avoiding another mistake: error and posterror neural activity associated with adaptive posterror behavior change. *Cognitive, Affective & Behavioral Neuroscience*, 7(4), 317–26. <https://doi.org/10.3758/CABN.7.4.317>

Hester, R., Madeley, J., Murphy, K., & Mattingley, J. B. (2009). Learning from errors: error-related neural activity predicts improvements in future inhibitory control performance. *Journal of Neuroscience*, 29(22), 7158–7165. <https://doi.org/10.1523/JNEUROSCI.4337-08.2009>

Hutchison, R. M., Womelsdorf, T., Allen, E. A., Bandettini, P. A., Calhoun, V. D., Corbetta, M., ... Chang, C. (2013). Dynamic functional connectivity: Promise, issues, and interpretations. *NeuroImage*, 80, 360–378. <https://doi.org/10.1016/j.neuroimage.2013.05.079>

Huys, Q. J. M., Maia, T. V., & Frank, M. J. (2016). Computational psychiatry as a bridge between neuroscience and clinical applications. *Nature Neuroscience*, 19(3), 404–413. <https://doi.org/10.1038/nn.4238>

Jentzsch, I., & Dudschg, C. (2009). Why do we slow down after an error? Mechanisms underlying the effects of posterror slowing. *Quarterly Journal of Experimental Psychology* (2006), 62(2), 209–18. <https://doi.org/10.1080/>

17470210802240655

Jocham, G., Klein, T. A., & Ullsperger, M. (2011). Dopamine-mediated reinforcement learning signals in the striatum and ventromedial prefrontal cortex underlie value-based choices. *Journal of Neuroscience*, 31(5), 1606–13. <https://doi.org/10.1523/JNEUROSCI.3904-10.2011>

Kahnt, T., Park, S. Q., Cohen, M. X., Beck, A., Heinz, A., & Wräse, J. (2009). Dorsal striatal-midbrain connectivity in humans predicts how reinforcements are used to guide decisions. *Journal of Cognitive Neuroscience*, 21(7), 1332–45. <https://doi.org/10.1162/jocn.2009.21092>

Kamin, L. J. (1969). Predictability, surprise, attention, and conditioning. In B. A. Campbell & M. R. Church (Eds.), *Punishment and aversive behavior* (pp. 279–296). Appleton-Century-Crofts.

Kerns, J. G., Cohen, J. D., MacDonald, A. W., Cho, R. Y., Stenger, V. A., & Carter, C. S. (2004). Anterior cingulate conflict monitoring and adjustments in control. *Science*, 303(5660), 1023–6. <https://doi.org/10.1126/science.1089910>

King, J. A., Korb, F. M., Cramon, D. Y. von, & Ullsperger, M. (2010). Post-error behavioral adjustments are facilitated by activation and suppression of task-relevant and task-irrelevant information processing. *Journal of Neuroscience*, 30(38), 12759–69. <https://doi.org/10.1523/JNEUROSCI.3274-10.2010>

Klein, T. A., Neumann, J., Reuter, M., Hennig, J., Von Cramon, D., & Ullsperger, M. (2007). Genetically determined differences in learning from errors. *Science*, 318(5856), 1642–1645. <https://doi.org/10.1126/science.1145044>

Kumar, P., Waiter, G., Ahearn, T., Milders, M., Reid, I., & Steele, J. D. (2008). Abnormal temporal difference reward-learning signals in major depression. *Brain*,

131(8), 2084–2093. <https://doi.org/10.1093/brain/awn136>

Laming, D. (1979). Choice reaction performance following an error. *Acta Psychologica*, 43(3), 199–224. [https://doi.org/10.1016/0001-6918\(79\)90026-X](https://doi.org/10.1016/0001-6918(79)90026-X)

Lee, S. W., Shimojo, S., & O'Doherty, J. P. (2014). Neural computations underlying arbitration between model-based and model-free learning. *Neuron*, 81(3), 687–99. <https://doi.org/10.1016/j.neuron.2013.11.028>

Lenartowicz, A., Kalar, D. J., Congdon, E., & Poldrack, R. A. (2010). Towards an Ontology of Cognitive Control. *Topics in Cognitive Science*, 2(4), 678–692. <https://doi.org/10.1111/j.1756-8765.2010.01100.x>

Logothetis, N. K., Pauls, J., Augath, M., Trinath, T., & Oeltermann, A. (2001). Neurophysiological investigation of the basis of the fMRI signal. *Nature*, 412(6843), 150–157. <https://doi.org/10.1038/35084005>

Madsen, K. S., Baaré, W. F. C., Vestergaard, M., Skimminge, A., Ejersbo, L. R., Ramsøy, T. Z., ... Jernigan, T. L. (2010). Response inhibition is associated with white matter microstructure in children. *Neuropsychologia*, 48(4), 854–862. <https://doi.org/10.1016/j.neuropsychologia.2009.11.001>

Maia, T. V., & Frank, M. J. (2011). From reinforcement learning models to psychiatric and neurological disorders. *Nature Neuroscience*, 14(2), 154–162. <https://doi.org/10.1038/nn.2723>

Makris, N., Biederman, J., Valera, E. M., Bush, G., Kaiser, J., Kennedy, D. N., ... Seidman, L. J. (2007). Cortical thinning of the attention and executive function networks in adults with attention-deficit/hyperactivity disorder. *Cerebral Cortex*, 17(6), 1364–1375. <https://doi.org/10.1093/cercor/bhl047>

Mathias, S. R., Knowles, E. E. M., Barrett, J., Leach, O., Buccheri, S., Beetham, T., ... Glahn, D. C. (2017). The Processing-Speed Impairment in Psychosis

Is More Than Just Accelerated Aging. *Schizophrenia Bulletin*, 43(4), 814–823. <https://doi.org/10.1093/schbul/sbw168>

Miller, E. K. (2000). The prefrontal cortex and cognitive control. *Nature Reviews Neuroscience*, 1(1), 59–65. <https://doi.org/10.1038/35036228>

Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529–533. <https://doi.org/10.1038/nature14236>

Montague, P. R., Dayan, P., & Sejnowski, T. J. (1996). A framework for mesencephalic dopamine systems based on predictive Hebbian learning. *Journal of Neuroscience*, 16(5), 1936–1947.

Montague, P. R., Dolan, R. J., Friston, K. J., & Dayan, P. (2012). Computational psychiatry. *Trends in Cognitive Sciences*, 16(1), 72–80. <https://doi.org/10.1016/j.tics.2011.11.018>

Moustafa, A. A., Kéri, S., Somlai, Z., Balsdon, T., Frydecka, D., Misiak, B., & White, C. (2015). Drift diffusion model of reward and punishment learning in schizophrenia: Modeling and experimental data. *Behavioural Brain Research*, 291, 147–154. <https://doi.org/10.1016/j.bbr.2015.05.024>

Niv, Y. (2009). Reinforcement learning in the brain. *Journal of Mathematical Psychology*, 53(3), 139–154. <https://doi.org/10.1016/j.jmp.2008.12.005>

Niv, Y., & Schoenbaum, G. (2008). Dialogues on prediction errors. *Trends in Cognitive Sciences*, 12(7), 265–72. <https://doi.org/10.1016/j.tics.2008.03.006>

Niv, Y., Edlund, J. A., Dayan, P., & O'Doherty, J. P. (2012). Neural prediction errors reveal a risk-sensitive reinforcement-learning process in the human brain. *Journal of Neuroscience*, 32(2), 551–562. <https://doi.org/10.1523/JNEUROSCI.1111-11.2012>

5498-10.2012

Notebaert, W., Houtman, F., Opstal, F. V., Gevers, W., Fias, W., & Verguts, T. (2009). Post-error slowing: an orienting account. *Cognition*, 111(2), 275–9. <https://doi.org/10.1016/j.cognition.2009.02.002>

Ogawa, S., Lee, T.-M., Kay, A. R., & Tank, D. W. (1990). Brain magnetic resonance imaging with contrast dependent on blood oxygenation. *Proceedings of the National Academy of Sciences*, 87(24), 9868–9872. <https://doi.org/10.1073/pnas.87.24.9868>

Otto, A. R., Skatova, A., Madlon-Kay, S., & Daw, N. D. (2015). Cognitive control predicts use of model-based reinforcement learning. *Journal of Cognitive Neuroscience*, 27(2), 319–33. https://doi.org/10.1162/jocn_a_00709

Pavlov, I. (1927). *Conditioned reflexes: an investigation of the physiological activity of the cerebral cortex*. London: Oxford Univ. Press.

Purcell, B. A., & Kiani, R. (2016). Neural Mechanisms of Post-error Adjustments of Decision Policy in Parietal Cortex. *Neuron*, 89(3), 658–671. <https://doi.org/10.1016/j.neuron.2015.12.027>

Rae, C. L., Hughes, L. E., Anderson, M. C., & Rowe, J. B. (2015). The Prefrontal Cortex Achieves Inhibitory Control by Facilitating Subcortical Motor Pathway Connectivity. *Journal of Neuroscience*, 35(2), 786–794. <https://doi.org/10.1523/JNEUROSCI.3093-13.2015>

Ratcliff, R., & Childers, R. (2015). Individual differences and fitting methods for the two-choice diffusion model of decision making. *Decision*, 2(4), 237–279. <https://doi.org/10.1037/dec0000030>

Ratcliff, R., Smith, P. L., Brown, S. D., & McKoon, G. (2016). Diffusion Decision Model: Current Issues and History. *Trends in Cognitive Sciences*, 20(4), 260–281.

<https://doi.org/10.1016/j.tics.2016.01.007>

Ratcliff, R., Thapar, A., & McKoon, G. (2010). Individual differences, aging, and IQ in two-choice tasks. *Cognitive Psychology*, 60(3), 127–157. <https://doi.org/10.1016/j.cogpsych.2009.09.001>

Redgrave, P., Gurney, K., & Reynolds, J. (2008). What is reinforced by phasic dopamine signals? *Brain Research Reviews*, 58(2), 322–39. <https://doi.org/10.1016/j.brainresrev.2007.10.007>

Redgrave, P., Prescott, T. J., & Gurney, K. (1999). Is the short latency dopamine burst too short to signal reinforcement error? *Trends in Neuroscience*, 22(4), 146–151. [https://doi.org/10.1016/S0166-2236\(98\)01373-3](https://doi.org/10.1016/S0166-2236(98)01373-3)

Rescorla, R. A., & Wagner, A. R. (1972). A theory of Pavlovian conditioning: Variations in the effectiveness of reinforcement and nonreinforcement. In A. H. Black & W. F. Prokasy (Eds.), *Classical conditioning: Current research and theory* (pp. 64–99). New York: Appleton Century Crofts.

Ridderinkhof, K. R., Wildenberg, W. P. M. van den, Segalowitz, S. J., & Carter, C. S. (2004). Neurocognitive mechanisms of cognitive control: the role of prefrontal cortex in action selection, response inhibition, performance monitoring, and reward-based learning. *Brain and Cognition*, 56(2), 129–40. <https://doi.org/10.1016/j.bandc.2004.09.016>

Roiser, J. P., Stephan, K. E., Ouden, H. E. M. den, Barnes, T. R. E., Friston, K. J., & Joyce, E. M. (2009). Do patients with schizophrenia exhibit aberrant salience? *Psychological Medicine*, 39(02), 199. <https://doi.org/10.1017/S0033291708003863>

Sabb, F. W., Bearden, C. E., Glahn, D. C., Parker, D. S., Freimer, N., & Bilder, R. M. (2008). A collaborative knowledge base for cognitive phenomics. *Molecular*

Psychiatry, 13(4), 350–60. <https://doi.org/10.1038/sj.mp.4002124>

Schiffer, A.-M., Muller, T., Yeung, N., & Waszak, F. (2014). Reward Activates Stimulus-Specific and Task-Dependent Representations in Visual Association Cortices. *Journal of Neuroscience*, 34(47), 15610–15620. <https://doi.org/10.1523/JNEUROSCI.1640-14.2014>

Schiffler, B. C., Almeida, R., Granqvist, M., & Bengtsson, S. L. (2016). Memory-reliant Posterror Slowing Is Associated with Successful Learning and Fronto-occipital Activity. *Journal of Cognitive Neuroscience*, 28(10), 1539–1552. https://doi.org/10.1162/jocn_a_00987

Schiffler, B. C., Bengtsson, S. L., & Lundqvist, D. (2017). The Sustained Influence of an Error on Future Decision-Making. *Frontiers in Psychology*, 8, 1077. <https://doi.org/10.3389/fpsyg.2017.01077>

Schultz, W., Dayan, P., & Montague, P. R. (1997). A Neural Substrate of Prediction and Reward. *Science*, 275(5306), 1593–1599. <https://doi.org/10.1126/science.275.5306.1593>

Seymour, B., O'Doherty, J. P., Dayan, P., Koltzenburg, M., Jones, A. K., Dolan, R. J., ... Frackowiak, R. S. (2004). Temporal difference models describe higher-order learning in humans. *Nature*, 429(6992), 664–667. <https://doi.org/10.1038/nature02581>

Siegert, S., Herrojo Ruiz, M., Brücke, C., Huebl, J., Schneider, G. H., Ullsperger, M., & Kühn, A. A. (2014). Error signals in the subthalamic nucleus are related to post-error slowing in patients with Parkinson's disease. *Cortex*, 60, 103–120. <https://doi.org/10.1016/j.cortex.2013.12.008>

Silver, D., Schrittwieser, J., Simonyan, K., Antonoglou, I., Huang, A., Guez, A., ... Hassabis, D. (2017). Mastering the game of Go without human knowledge.

Nature, 550(7676), 354–359. <https://doi.org/10.1038/nature24270>

Skinner, B. F. (1935). The Generic Nature of the Concepts of Stimulus and Response. *The Journal of General Psychology*, 12(1), 40–65. <https://doi.org/10.1080/00221309.1935.9920087>

Smittenaar, P., FitzGerald, T. H. B., Romei, V., Wright, N. D., & Dolan, R. J. (2013). Disruption of Dorsolateral Prefrontal Cortex Decreases Model-Based in Favor of Model-free Control in Humans. *Neuron*, 80(4), 914–919. <https://doi.org/10.1016/j.neuron.2013.08.009>

Steinbauer, M., Ernst, B., & I bald, K. W. (2017). Isolating component processes of posterror slowing with the psychological refractory period paradigm. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 43(4), 653–659. <https://doi.org/10.1037/xlm0000329>

Storsve, A. B., Fjell, A. M., Tamnes, C. K., Westlye, L. T., Overbye, K., Aasland, H. W., & Walhovd, K. B. (2014). Differential Longitudinal Changes in Cortical Thickness, Surface Area and Volume across the Adult Life Span: Regions of Accelerating and Decelerating Change. *Journal of Neuroscience*, 34(25), 8488–8498. <https://doi.org/10.1523/JNEUROSCI.0391-14.2014>

Sutton, R. S. (1988). Learning to Predict by the Methods of Temporal Differences. *Machine Learning*, 3(1), 9–44. <https://doi.org/10.1007/BF00115009>

Sutton, R. S., & Barto, A. G. (1990). Time-Derivative Models of Pavlovian Reinforcement. In M. Gabriel & J. Moore (Eds.), *Learning and computational neuroscience: Foundations of adaptive networks* (pp. 497–537). Cambridge: MIT Press. <https://doi.org/10.1017/CBO9780511541354.012>

Sutton, R. S., & Barto, A. G. (1998). *Reinforcement Learning: An Introduction*.

Cambridge: MIT Press.

Tabibnia, G., Monterosso, J. R., Baicy, K., Aron, A. R., Poldrack, R. A., Chakrapani, S., ... London, E. D. (2011). Different Forms of Self-Control Share a Neurocognitive Substrate. *Journal of Neuroscience*, 31(13), 4805–4810. <https://doi.org/10.1523/JNEUROSCI.2859-10.2011>

Thompson, W. H., Brantefors, P., & Fransson, P. (2017). From static to temporal network theory - applications to functional brain connectivity. *Network Neuroscience*, 1(2), 69–99. https://doi.org/10.1162/netn_a_00011

Thorndike, E. L. (1911). *Animal intelligence: Experimental studies*. Macmillan.

Ullsperger, M., & Danielmeier, C. (2016). Reducing Speed and Sight: How Adaptive Is Post-Error Slowing? *Neuron*, 89(3), 430–432. <https://doi.org/10.1016/j.neuron.2016.01.035>

Ullsperger, M., Danielmeier, C., & Jocham, G. (2014). Neurophysiology of performance monitoring and adaptive behavior. *Physiological Reviews*, 94(1), 35–79. <https://doi.org/10.1152/physrev.00041.2012>

Van Den Bos, W., Cohen, M. X., Kahnt, T., & Crone, E. A. (2012). Striatum-medial prefrontal cortex connectivity predicts developmental changes in reinforcement learning. *Cerebral Cortex*, 22(6), 1247–1255. <https://doi.org/10.1093/cercor/bhr198>

Volkow, N. D., Fowler, J. S., Wang, G. J., Telang, F., Logan, J., Jayne, M., ... Swanson, J. M. (2010). Cognitive control of drug craving inhibits brain reward regions in cocaine abusers. *NeuroImage*, 49(3), 2536–2543. <https://doi.org/10.1016/j.neuroimage.2009.10.088>

Voon, V., Derbyshire, K., Rück, C., Irvine, M. A., Worbe, Y., Enander, J., ... Bullmore, E. T. (2015). Disorders of compuslivity: a common bias towards learning

- habits. *Molecular Psychiatry*, 20(3), 345–352. <https://doi.org/10.1038/mp.2014.44>
- Watkins, C. J. C. H. (1989). *Learning from Delayed Rewards* (Doctoral Dissertation). King's College, Cambridge.
- Watkins, C. J. C. H., & Dayan, P. (1992). Technical Note: Q-Learning. *Machine Learning*, 8(3), 279–292. <https://doi.org/10.1023/A:1022676722315>
- Wiecki, T. V., Sofer, I., & Frank, M. J. (2013). HDDM: Hierarchical Bayesian estimation of the Drift-Diffusion Model in Python. *Frontiers in Neuroinformatics*, 7. <https://doi.org/10.3389/fninf.2013.00014>
- Wunderlich, K., Smittenaar, P., & Dolan, R. J. (2012). Dopamine enhances model-based over model-free choice behavior. *Neuron*, 75(3), 418–24. <https://doi.org/10.1016/j.neuron.2012.03.042>
- Zaghoul, K. A., Weidemann, C. T., Lega, B. C., Jaggi, J. L., Baltuch, G. H., & Kahana, M. J. (2012). Neuronal Activity in the Human Subthalamic Nucleus Encodes Decision Conflict during Action Selection. *Journal of Neuroscience*, 32(7), 2453–2460. <https://doi.org/10.1523/JNEUROSCI.5815-11.2012>
- Zhang, J., Rittman, T., Nombela, C., Fois, A., Coyle-Gilchrist, I., Barker, R. A., ... Rowe, J. B. (2016). Different decision deficits impair response inhibition in progressive supranuclear palsy and Parkinson's disease. *Brain*, 139(1), 161–173. <https://doi.org/10.1093/brain/awv331>